

# **What is artificial intelligence?**

---

Artificial intelligence is the ability of a computer to perform tasks commonly associated with intelligent beings.

# What is machine learning?

---

Machine learning is the study of algorithms that learn from examples and experience instead of relying on hard-coded rules and make predictions on new data.

# What is deep learning?

---

Deep learning is a subfield of machine learning focusing on learning data representations as successive layers of increasingly meaningful representations.

# ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



1950's      1960's      1970's      1980's      1990's      2000's      2010's

## MACHINE LEARNING

Machine learning begins to flourish.



## DEEP LEARNING

Deep learning breakthroughs drive AI boom.

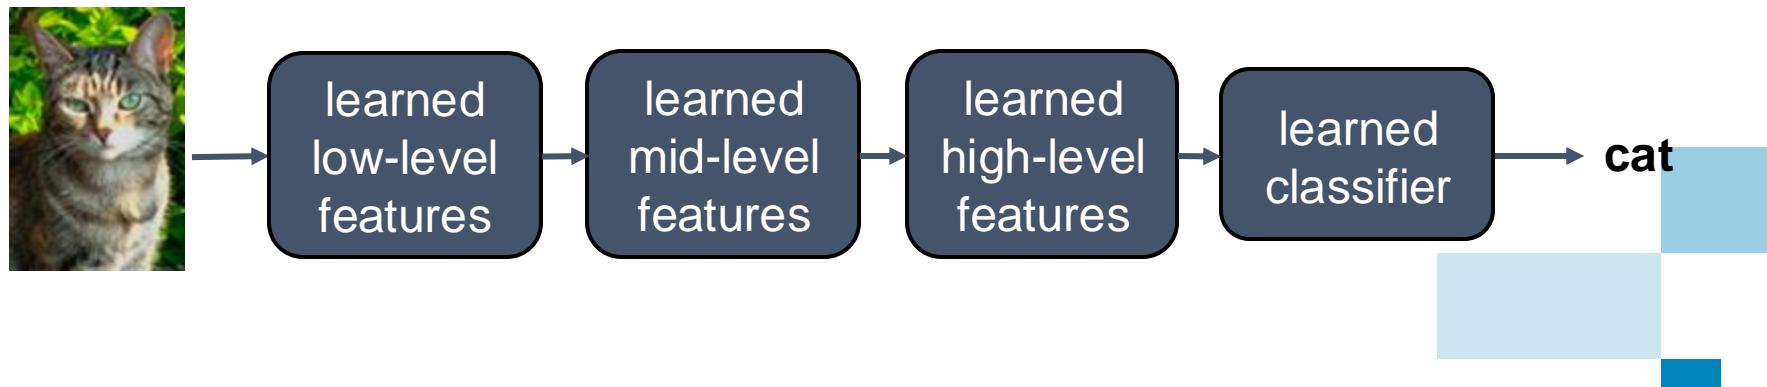


Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

## “Traditional” machine learning:

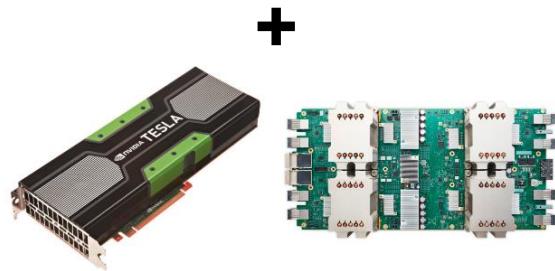


## Deep, “end-to-end” learning:



# Why Deep Learning Now?

- Big Data
  - Larger Datasets
  - Easier Collection and Storage
- Hardware
  - GPUs
  - Massively Parallelizable Training
- Software
  - Improved Techniques
  - New Models
  - Toolboxes



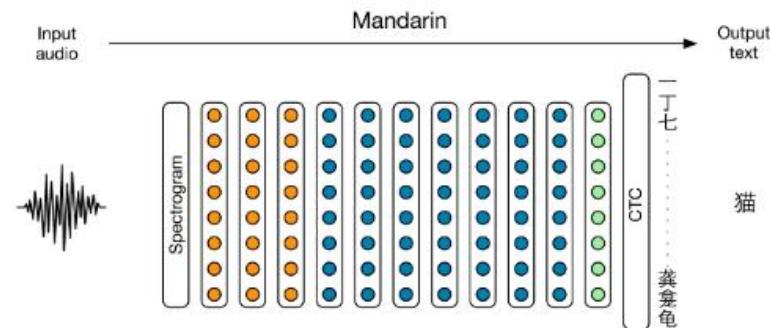
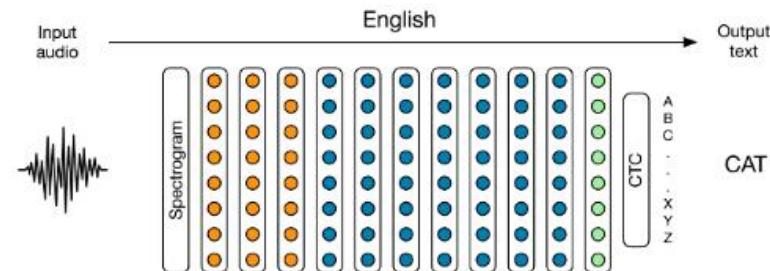
GPU and TPU





Deep  
Learning  
**BOOM!!!**

# DL Today: Speech-to-Text



- Convolution Layer
- Recurrent Layer
- Fully Connected Layer

[Baidu 2014]

# DL Today: Vision



[Krizhevsky 2012]



[Ciresan et al. 2013]



[Faster R-CNN - Ren 2015]

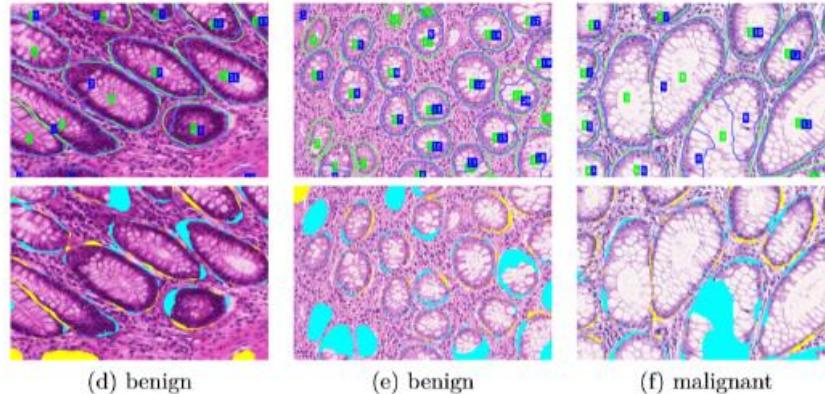


[NVIDIA dev blog]

# DL Today: Vision



[Stanford 2017]



[Nvidia Dev Blog 2017]

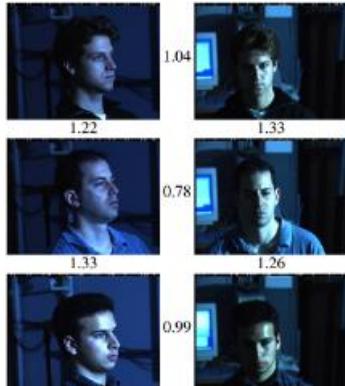
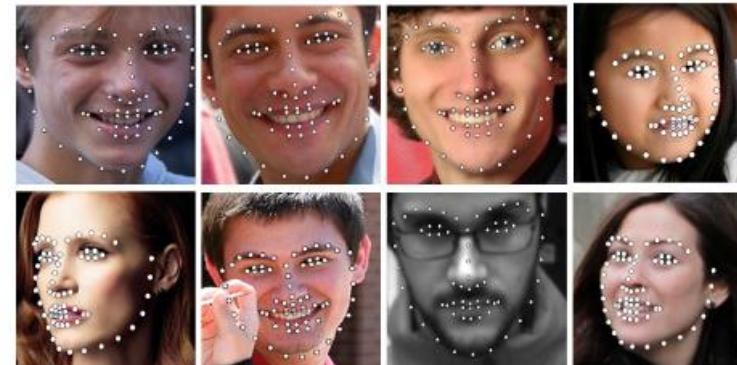


Figure 1. Illumination and Pose invariance.

[FaceNet - Google 2015]



[Facial landmark detection CUHK 2014]

# DL Today: NLP



Salit Kulla

to me

11:29 AM \*\*\*

Hey, Wynton Marsalis is playing this weekend. Do you have a preference between Saturday and Sunday?

-S

I'm down for either.

Let's do Saturday.

I'm fine with whatever.



Reply



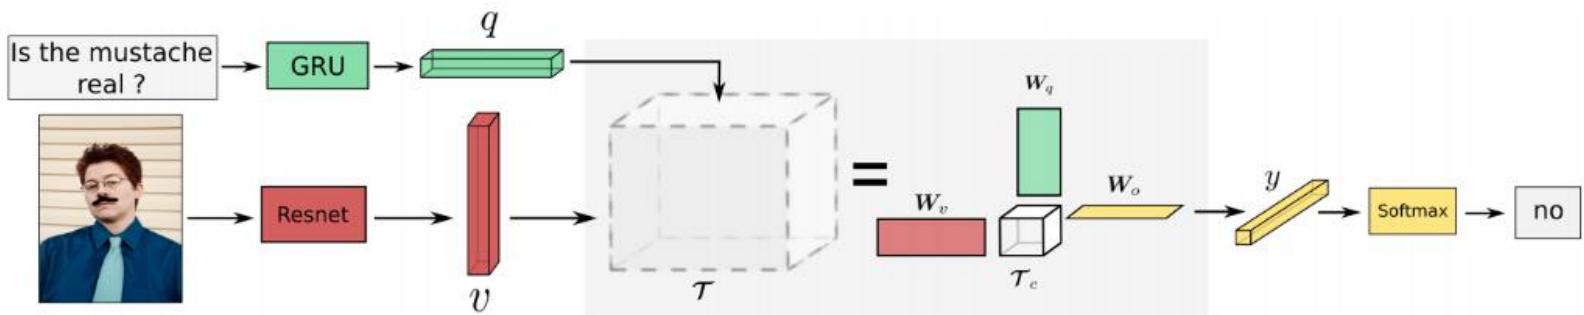
Forward



[Google Inbox Smart Reply]

[Amazon Echo / Alexa]

# DL Today: Vision + NLP



[VQA - Mutan 2017]



"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."



"boy is doing backflip on wakeboard."

[Karpathy 2015]

# DL Today: Generative models



Sampled celebrities [Nvidia 2017]

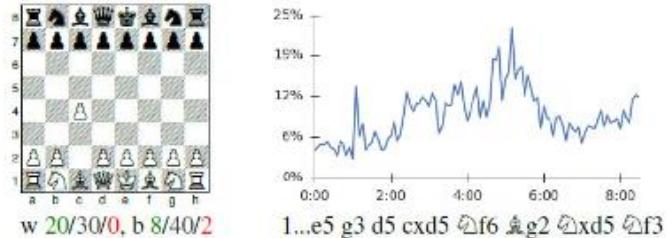
Text description	This bird is blue with white and has a very short beak	This bird has wings that are brown and has a yellow belly	A white bird with a black crown and yellow beak	This bird is white, black, and brown in color, with a brown beak	The bird has small beak, with reddish brown crown and gray belly	This is a small, black bird with a white breast and white on the wingbars.	This bird is white black and yellow in color, with a short black beak
Stage-I images							
Stage-II images							

StackGAN v2 [Zhang 2017]

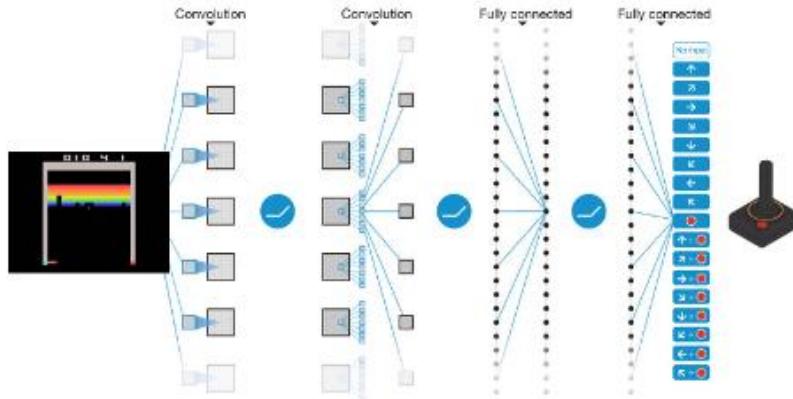
# DL for AI in games



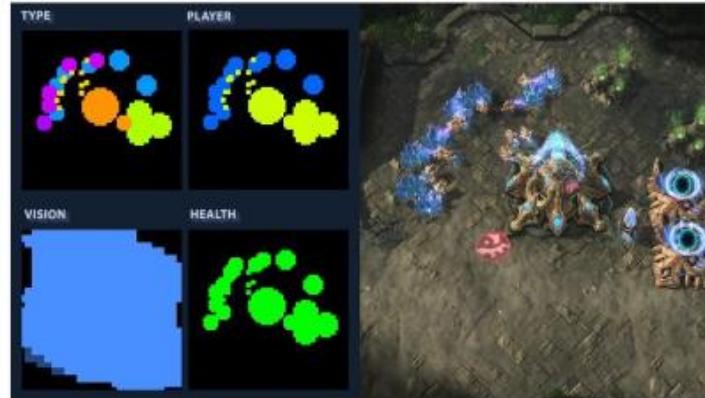
A10: English Opening



[Deepmind AlphaGo / Zero 2017]

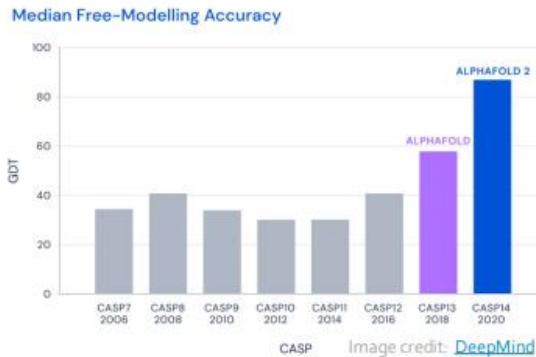


[Atari Games - DeepMind 2016]



[Starcraft 2 for AI research]

# AlphaFold – Protein Structure Prediction



**AlphaFold's AI could change the world of biological science as we know it**

DeepMind's latest AI breakthrough can accurately predict the way proteins fold

Has Artificial Intelligence 'Solved' Biology's Protein-Folding Problem?

12-14-20

**DeepMind's latest AI breakthrough could turbocharge drug discovery**

# ChatGPT - Large Language Models

## Microsoft Confirms Its \$10 Billion Investment Into ChatGPT, Changing How Microsoft Competes With Google, Apple And Other Tech Giants

generate python code to add two arrays and then reverse the second half of resultant array

```
python
import numpy as np

# Define two arrays
array1 = np.array([1, 2, 3, 4])
array2 = np.array([5, 6, 7, 8])

# Add the two arrays
result = array1 + array2

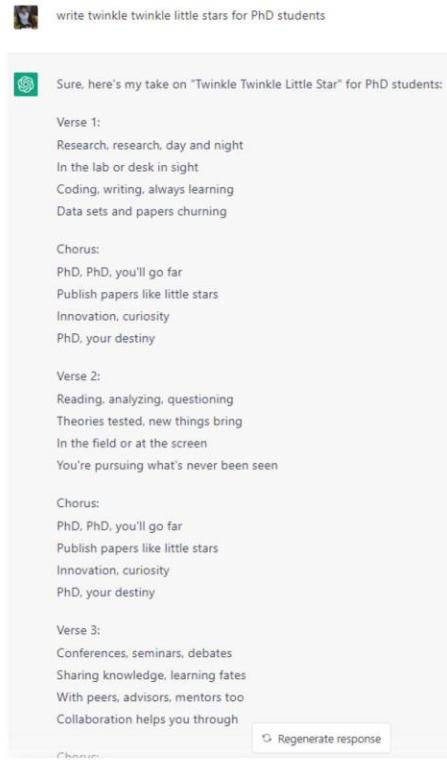
# Reverse the second half of the result array
mid = len(result) // 2
result[mid:] = result[mid:][::-1]

# Print the final result
print(result)
```

World

**ChatGPT breaks record with 100 million users – and investors come flocking**

Bianke Neethling • 15 Feb 2023



Google poised to release chatbot technology after ChatGPT success

# Stable Diffusion - Generative Models

A pikachu fine dining with a view to the Eiffel Tower

low quality

Generate image

The collage consists of four square images arranged in a 2x2 grid. Top-left: Pikachu sitting at a table with a plate of food, with the Eiffel Tower visible through a window. Top-right: A restaurant interior with yellow chairs and tables, and the Eiffel Tower seen through a glass partition. Bottom-left: Pikachu standing outdoors with its ears perked, with the Eiffel Tower in the background. Bottom-right: Pikachu sitting at a table on a rooftop terrace with the Eiffel Tower in the background.

"Super cute fluffy cat warrior in armor, photorealistic, 4K, ultra detailed, vray ren

high quality

Generate image

The collage consists of four square images arranged in a 2x2 grid. Top-left: A fluffy white and brown cat wearing a metallic suit of armor with shoulder guards and a chest plate. Top-right: A grey and white cat wearing a more ornate, golden-colored suit of armor with a plumed helmet. Bottom-left: A brown and white cat wearing a dark, segmented suit of armor. Bottom-right: A light-colored cat with large blue eyes wearing a detailed suit of armor with gold accents and a belt.

# 3D model generation

Instant Neural Graphics Primitives

Frame: 1.90 ms (525.9 FPS); Mem: 71.2 MB

► Training

▼ Rendering

Connect to VR/AR headset

Render : 0.8ms for 1920x1061 (4384 spp) VSync  
DLSS (Vulkan was missing at compilation time)

Dynamic resolution 20.0 Target FPS

Shade ▾ Render mode

Identity ▾ Tonemap curve

R: 0 G: 0 B: 0 A:255 Background

0.000 Exposure

World transform

Advanced rendering options

Debug visualization

Camera

First person controls Smooth motion Autofocus

0.000 Aperture size 0.000 Focus depth

50.625 Field of view 1.000 Zoom

Advanced camera settings

Snapshot

Snapshot Save Load Dump parameters as images w/ optimizer state  
base.ingp File Compress

Histograms of encoding parameters

Go to python REPL

Screen Recorder is sharing your screen. Stop sharing Hide

Type here to search

31°C Haze ENG 9:13 PM 24-Mar-23

# History of AI: Major Milestones

Table 1: Major milestones that will be covered in this paper

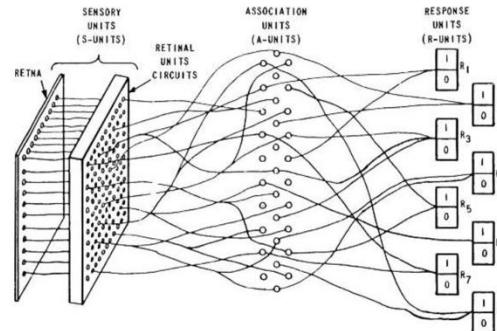
Year	Contributer	Contribution
300 BC	Aristotle	introduced Associationism, started the history of human's attempt to understand brain.
1873	Alexander Bain	introduced Neural Groupings as the earliest models of neural network, inspired Hebbian Learning Rule.
1943	McCulloch & Pitts	introduced MCP Model, which is considered as the ancestor of Artificial Neural Model.
1949	Donald Hebb	considered as the father of neural networks, introduced Hebbian Learning Rule, which lays the foundation of modern neural network.
1958	Frank Rosenblatt	introduced the first perceptron, which highly resembles modern perceptron.
1974	Paul Werbos	introduced Backpropagation
1980	Teuvo Kohonen Kunihiko Fukushima	introduced Self Organizing Map introduced Neocogitron, which inspired Convolutional Neural Network
1982	John Hopfield	introduced Hopfield Network
1985	Hilton & Sejnowski	introduced Boltzmann Machine
1986	Paul Smolensky Michael I. Jordan	introduced Harmonium, which is later known as Restricted Boltzmann Machine defined and introduced Recurrent Neural Network
1990	Yann LeCun	introduced LeNet, showed the possibility of deep neural networks in practice
1997	Schuster & Paliwal Hochreiter & Schmidhuber	introduced Bidirectional Recurrent Neural Network introduced LSTM, solved the problem of vanishing gradient in recurrent neural networks
2006	Geoffrey Hinton	introduced Deep Belief Networks, also introduced layer-wise pretraining technique, opened current deep learning era.
2009	Salakhutdinov & Hinton	introduced Deep Boltzmann Machines
2012	Geoffrey Hinton	introduced Dropout, an efficient way of training neural networks

# History of AI: Major Milestones



Dartmouth AI Conference 1956

*Early neural networks*



Perceptron, one of the first neural network architectures

Rosenblatt 1957



F. Rosenblatt

1957

The Perceptron, 1957

# History of AI: Major Milestones

## Early hype

"First serious rival to the human brain even devised."

"Remarkable machine capable of what amounts to thought"

— The New Yorker



1958

Manson, Stewart, Gill 1958

## Early hype

MASSACHUSETTS INSTITUTE OF TECHNOLOGY  
PROJECT MAC

Artificial Intelligence Group  
Vision Memo. No. 100.

July 7, 1966

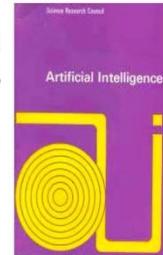
THE SUMMER VISION PROJECT  
Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

## *The Lighthill Report*

Papert 1966

"Most workers in AI research and in related fields confess to a pronounced feeling of disappointment in what has been achieved in the past twenty-five years. [...] In no part of the field have the discoveries made so far produced the major impact that was then promised."



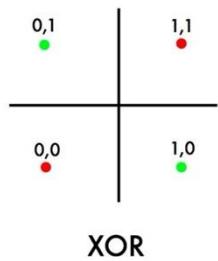
J. Lighthill

1972

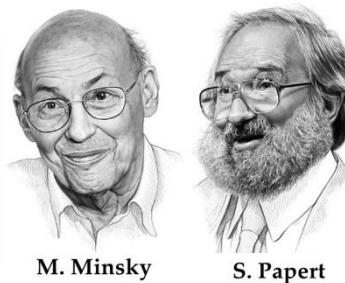
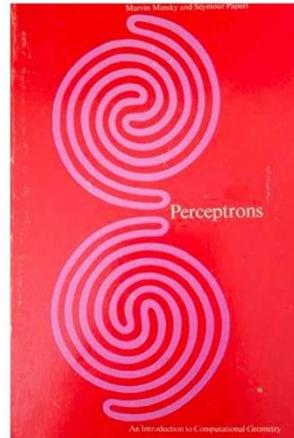
Lighthill 1972

# History of AI: Major Milestones

*The “XOR Affair”*



“[simple] perceptron  
cannot represent even  
the XOR function”



M. Minsky

S. Papert

1969

Minsky, Papert 1969

The XOR Affair, 1969

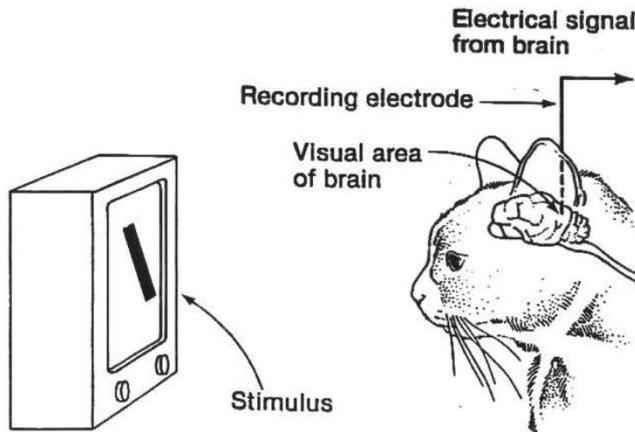
# History of AI: Major Milestones



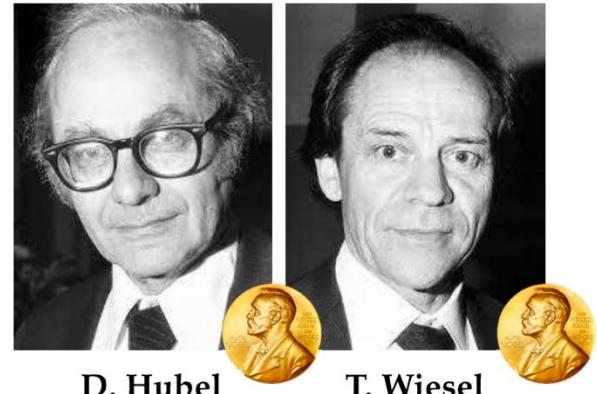
“AI WINTER”

# History of AI: Major Milestones

*Secrets of the visual cortex*



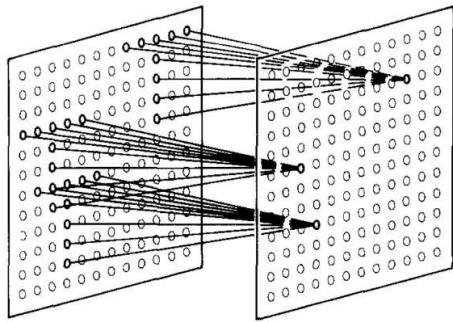
Experiments of Hubel and Wiesel that established the structure of the visual cortex



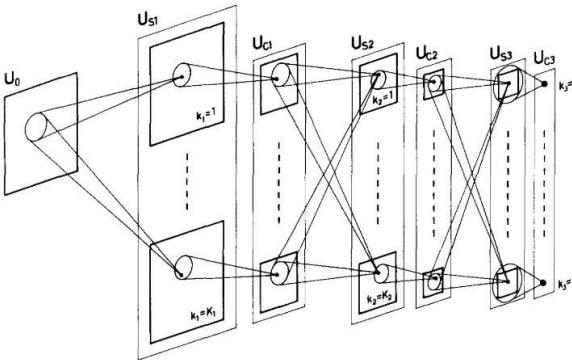
1959

# History of AI: Major Milestones

## *Neocognitron*



Neocognitron, an early geometric neural network



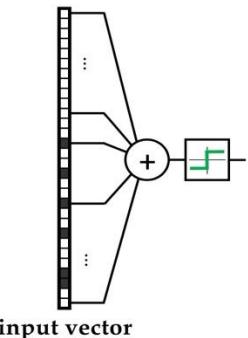
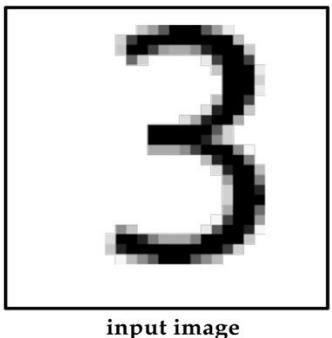
K. Fukushima

1980

Fukushima 1980

Neocognitron, 1980

# History of AI: Major Milestones



*Neocognitron*

- Deep neural network (7 layers tested)
- Local connectivity (“receptive fields”)
- Nonlinear filters with shared weights (S-layers)
- Average pooling (C-layers)
- ReLU activation function
- “Self-organised” (unsupervised) – no backprop yet!



K. Fukushima

1980

Fukushima 1980

Neocognitron, 1980

# History of AI: Major Milestones

*How to train your neural network?*



**F. Rosenblatt**

Perceptron  
learning rule  
(1 layer)



**A. Ivakhnenko**

Group method of  
data handling



**S. Linnainmaa**



**P. Werbos**

Backpropagation

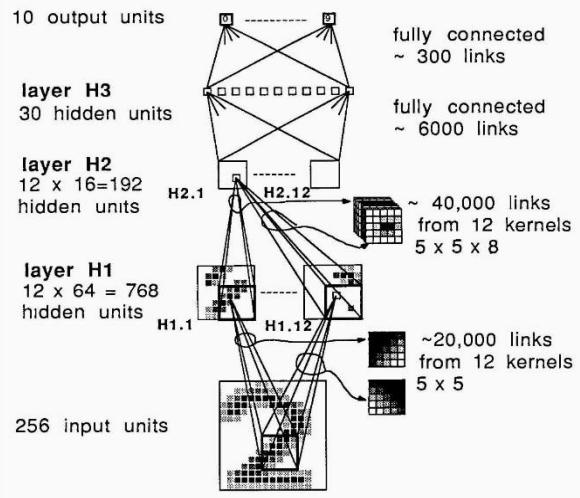


**D. Rumelhart**

Rosenblatt 1957; Ivakhnenko, Lapa 1966; Linnainmaa 1970; Werbos 1982; Rumelhart et al. 1986

# History of AI: Major Milestones

## *Convolutional neural networks*



First version of a CNN



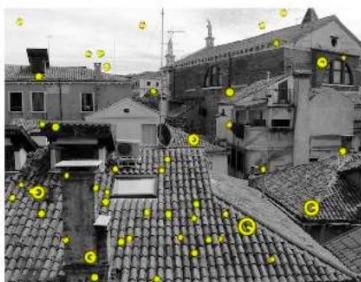
AT&T DSP-32C  
capable of 125m floating  
point multiply-accumulate  
operations/sec



Y. LeCun

# History of AI: Major Milestones

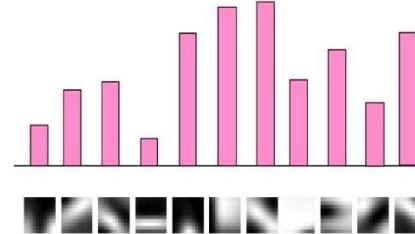
*Computer vision in the 2000s*



Feature detection



Feature description



Feature aggregation



Classification

A typical image classification pipeline from the 2000s

# Common Techniques for Feature Detection Before 2000

- **Harris Corner Detection** : Detects corners in an image by analyzing changes in intensity in different directions.

---

- **Canny Edge Detection:** Detects edges in an image by finding areas of rapid intensity change.
- **Scale-Invariant Feature Transform (SIFT):** Detects keypoints that are invariant to scale and rotation changes.

# Handcrafter features/kernel matrices



Original



Sharpen



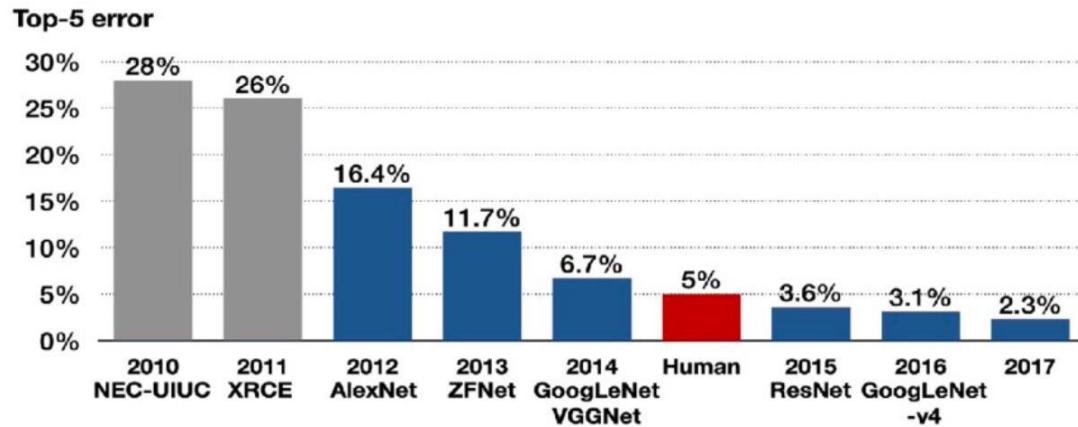
Edge Detect



“Strong” Edge  
Detect

# History of AI: Major Milestones

## *ImageNet*

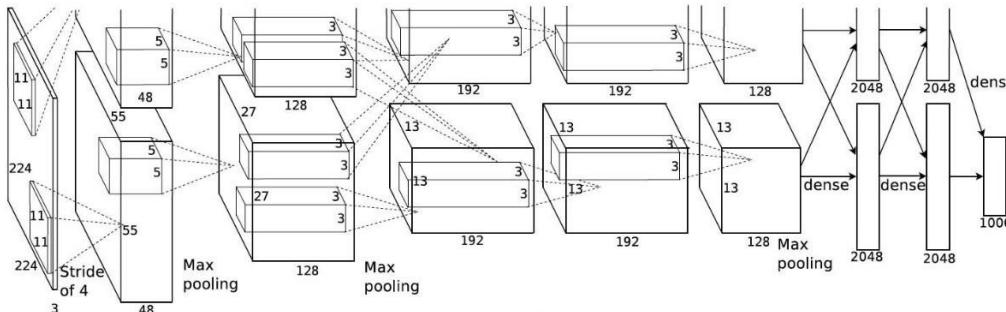


L. Fei-Fei

AlexNet beating all “handcrafted” approaches on ImageNet benchmark—the moment of truth for computer vision

# History of AI: Major Milestones

## AlexNet



AlexNet architecture

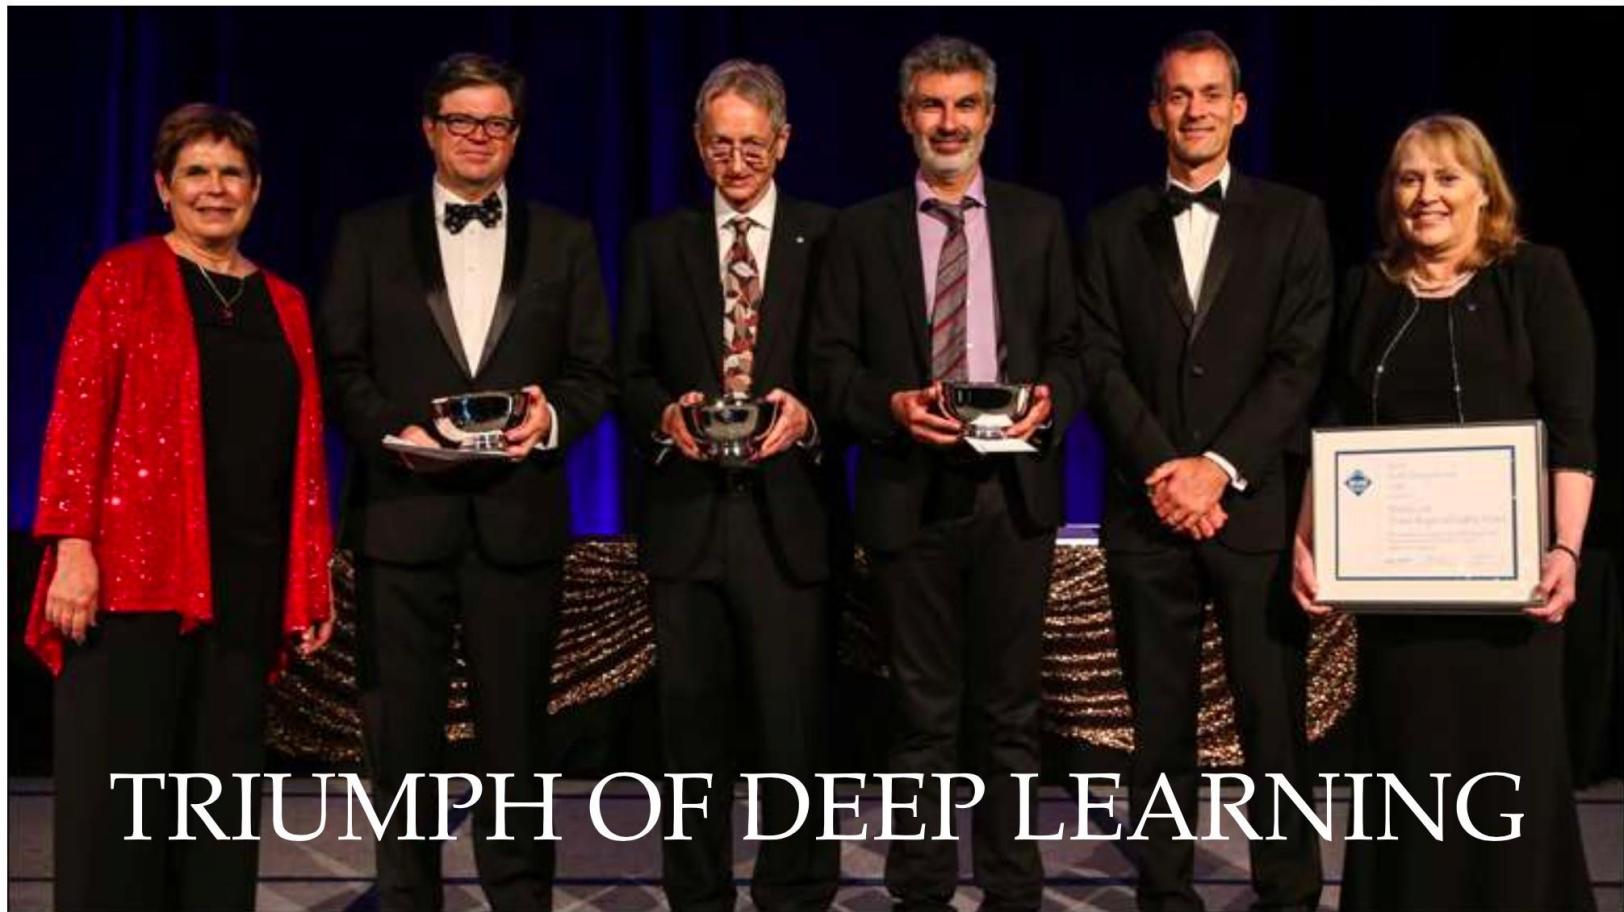
Nvidia GTX 580 GPU capable of  
~200G FLOP / sec



A. Krizhevsky

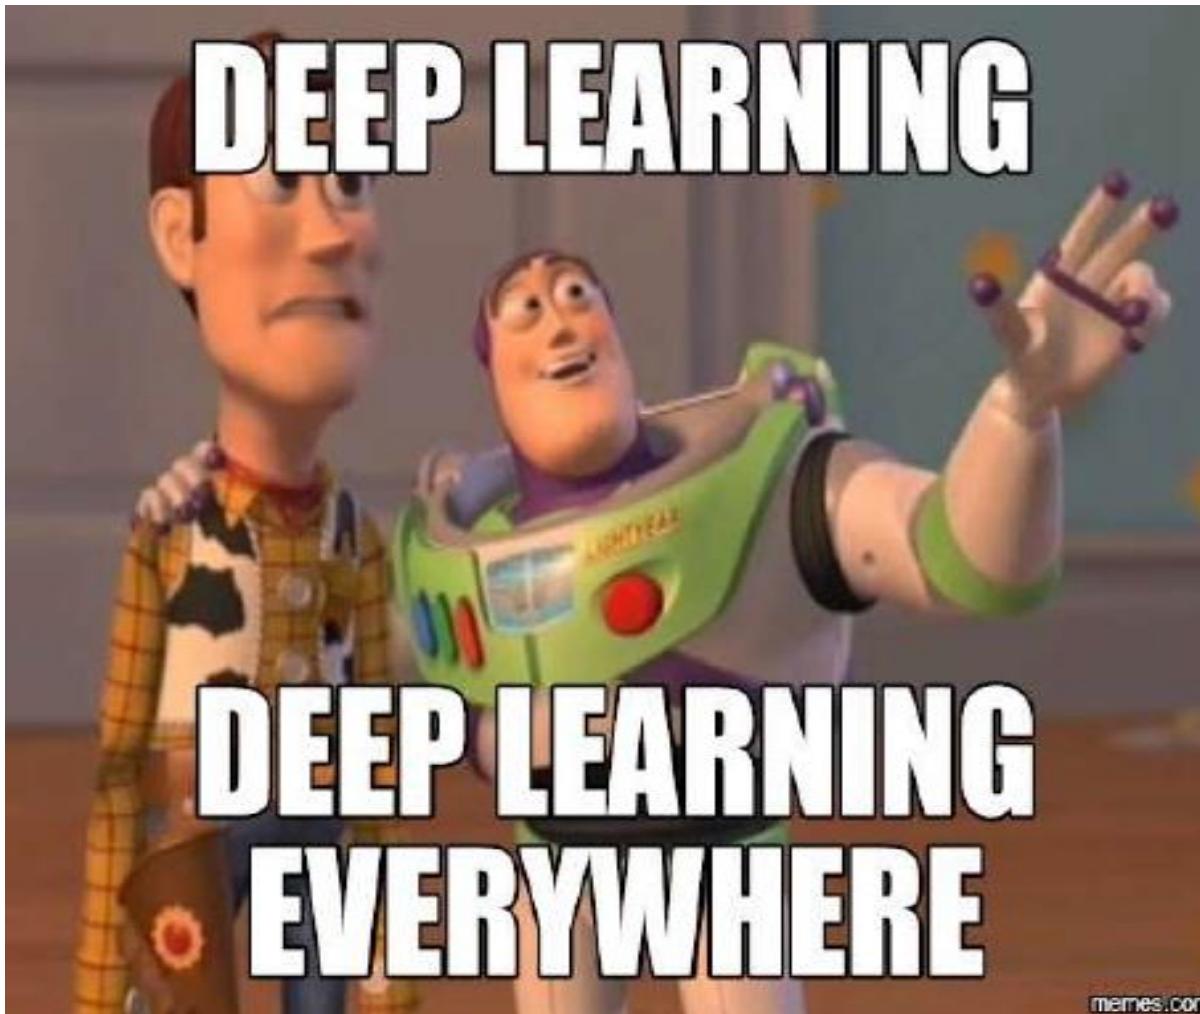
Krizhevsky et al. 2012

# History of AI: Major Milestones



TRIUMPH OF DEEP LEARNING

# Deep Learning Memes

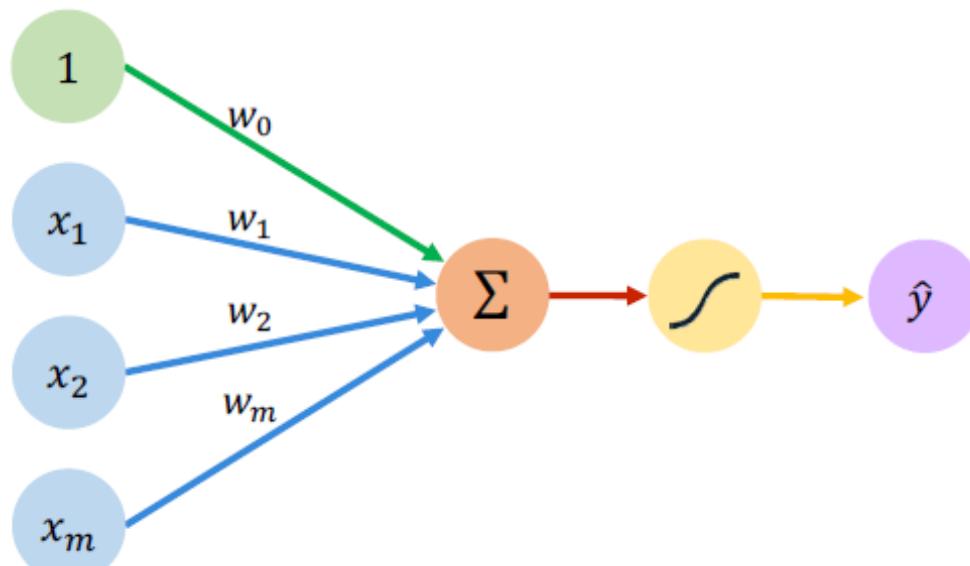


# The Perceptron

---

The structural building block of deep learning

# The Perceptron: Forward Propagation



Inputs

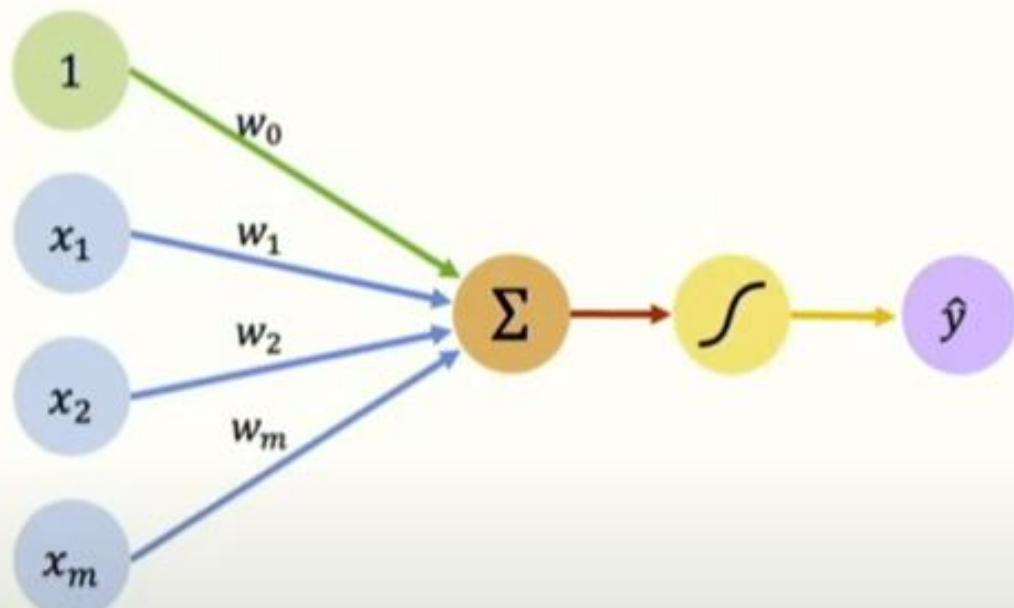
Weights

Sum

Non-Linearity

Output

# The Perceptron: Forward Propagation



Output

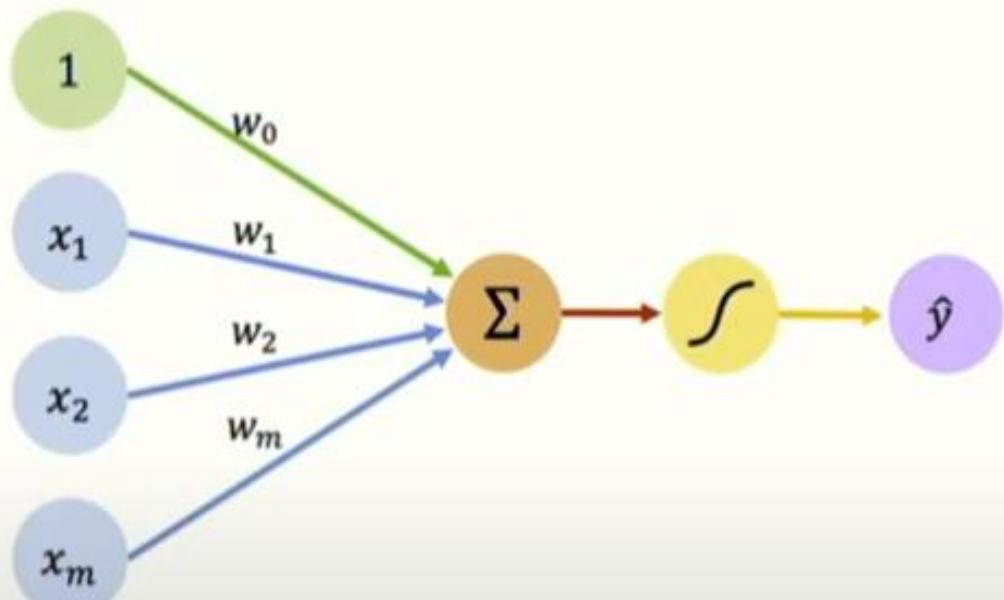
Linear combination of inputs

$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function

Bias

# The Perceptron: Forward Propagation



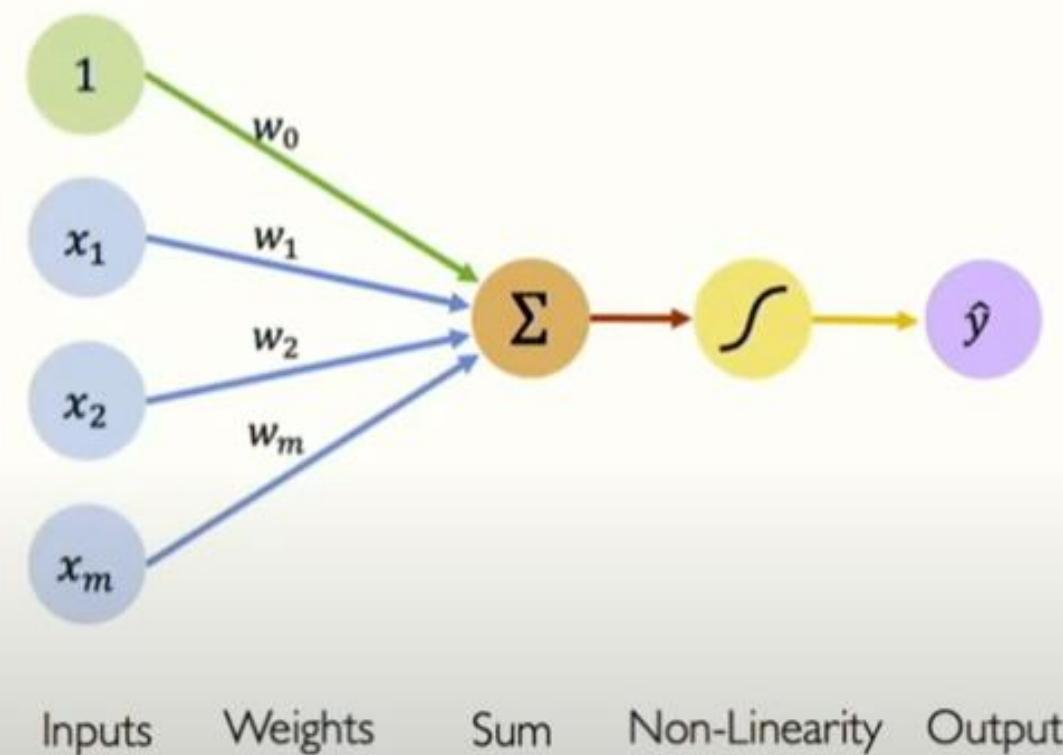
Inputs      Weights      Sum      Non-Linearity      Output

$$\hat{y} = g \left( w_0 + \sum_{l=1}^m x_l w_l \right)$$

$$\hat{y} = g ( w_0 + \mathbf{X}^T \mathbf{W} )$$

where:  $\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$  and  $\mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$

# The Perceptron: Forward Propagation

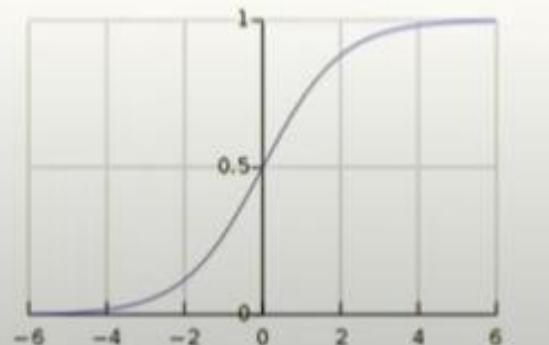


## Activation Functions

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$

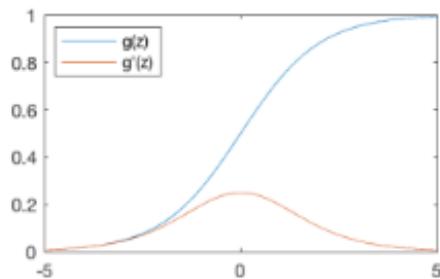
- Example: sigmoid function

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$



# Activation Functions

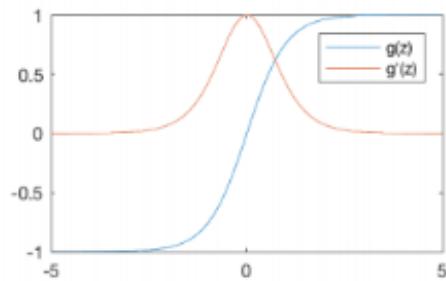
Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

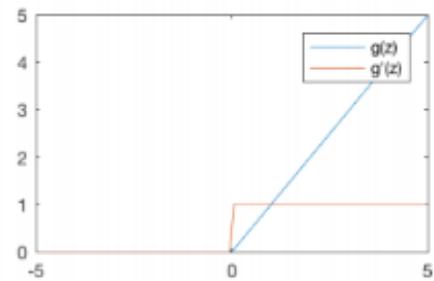
Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$g'(z) = 1 - g(z)^2$$

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

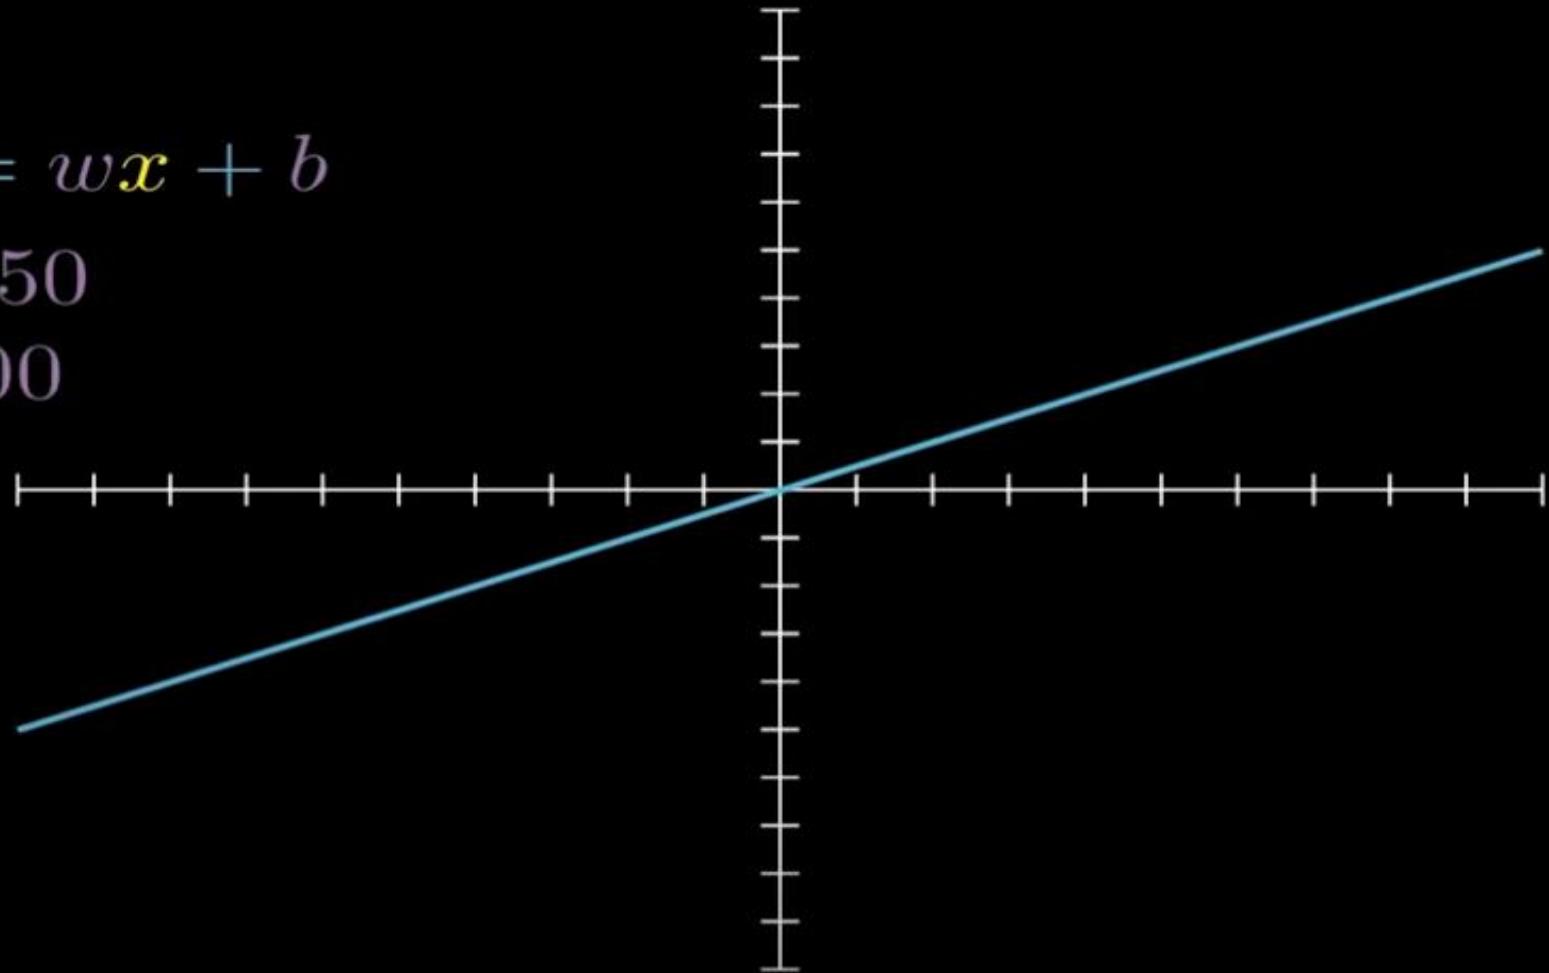
$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

# Neuron as a linear function

$$N(x) = wx + b$$

$$w = 0.50$$

$$b = 0.00$$



# Example: Estimating a target function using multiple neurons

$$N_0(x) = -5x + -7.7$$

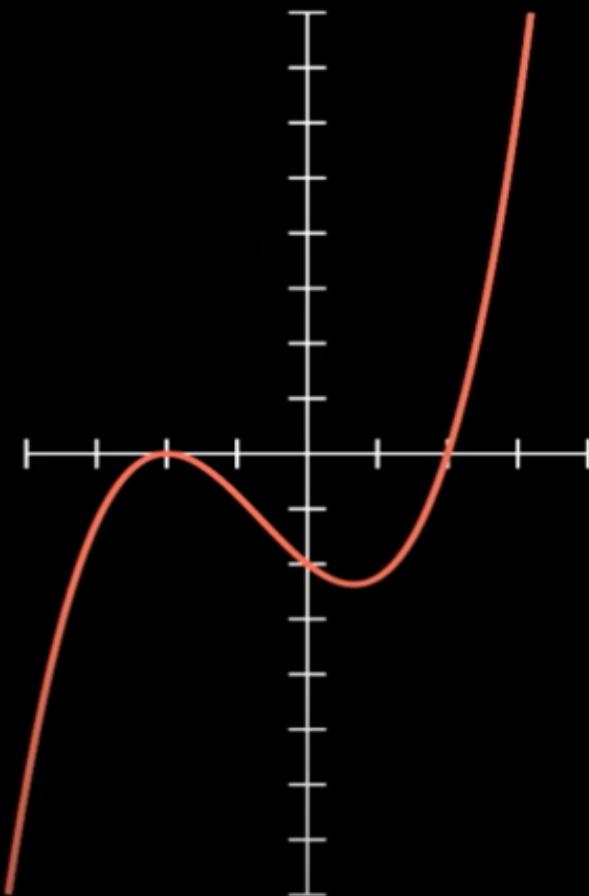
$$N_1(x) = -1.2x + -1.3$$

$$N_2(x) = 1.2x + 1$$

$$N_3(x) = 1.2x + -0.2$$

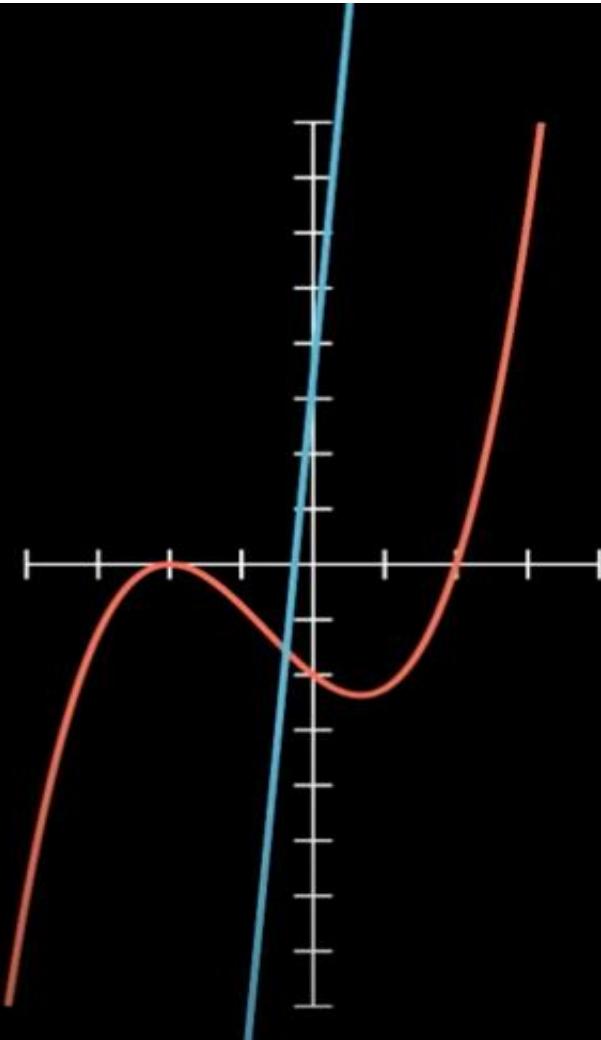
$$N_4(x) = 2x + -1.1$$

$$N_5(x) = 5x + -5$$



# No Non-linearity

$-1(-5x + -7.7) +$   
 $-1(-1.2x + -1.3) +$   
 $-1(1.2x + 1) +$   
 $1(1.2x + -0.2) +$   
 $1(2x + -1.1) +$   
 $1(5x + -5)$

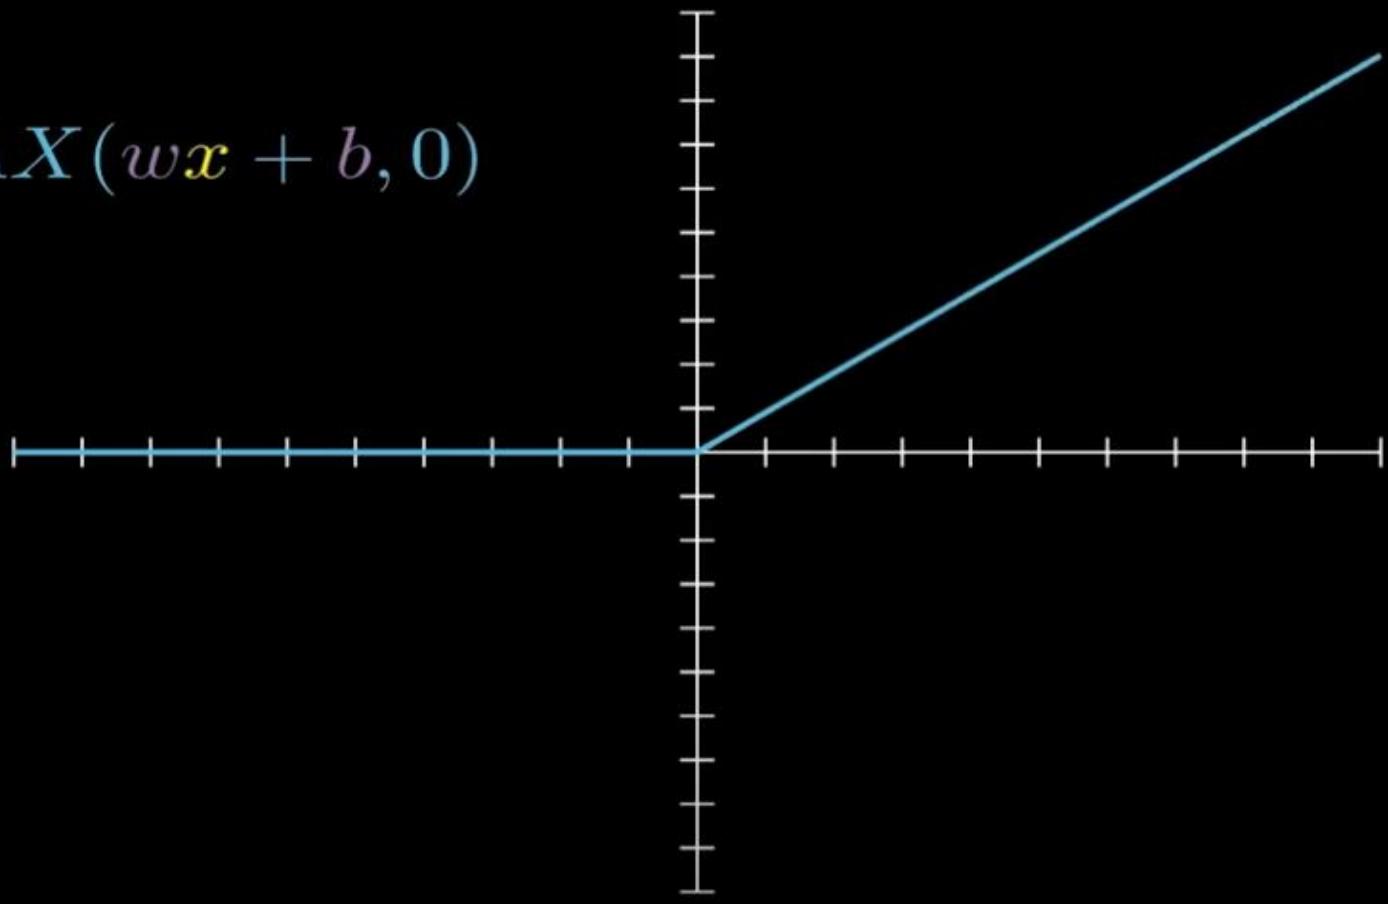


# Inducing Non-linearity using ReLU

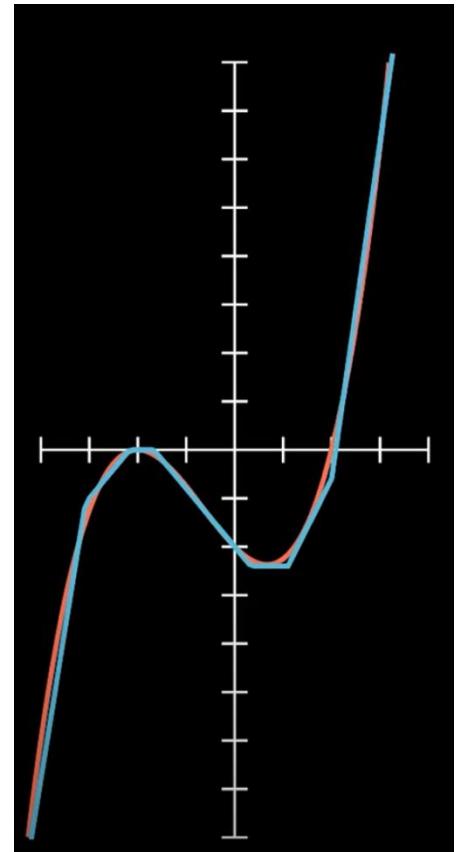
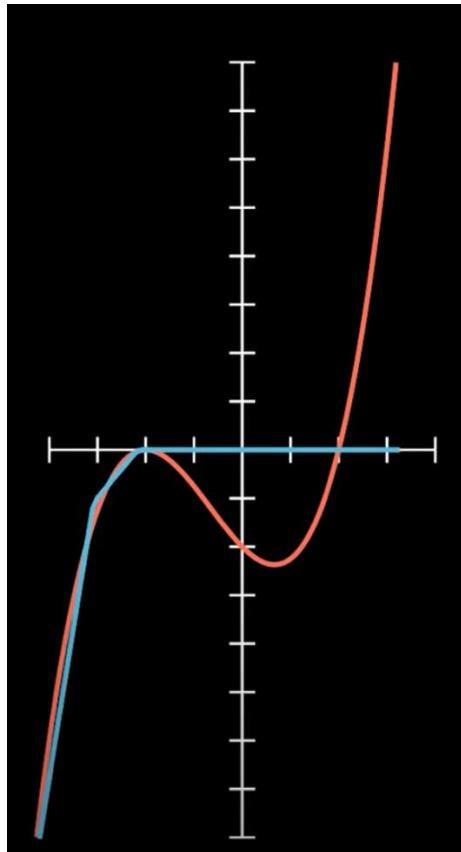
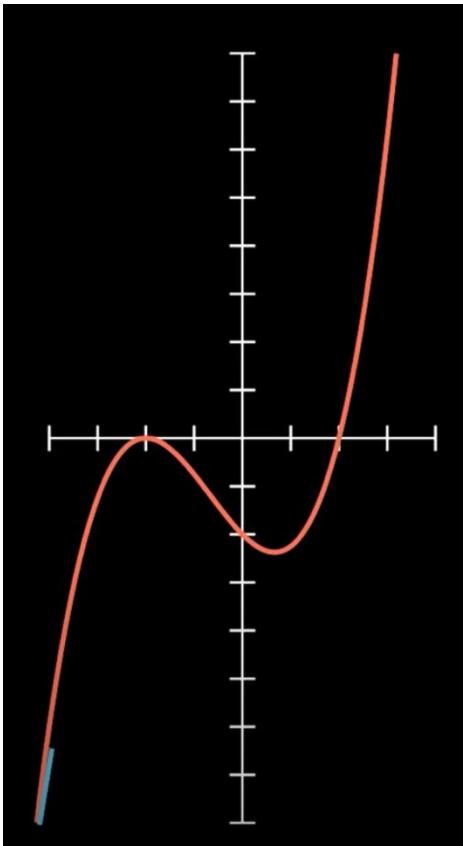
$$N(x) = \text{MAX}(wx + b, 0)$$

$$w = 0.90$$

$$b = 0.00$$

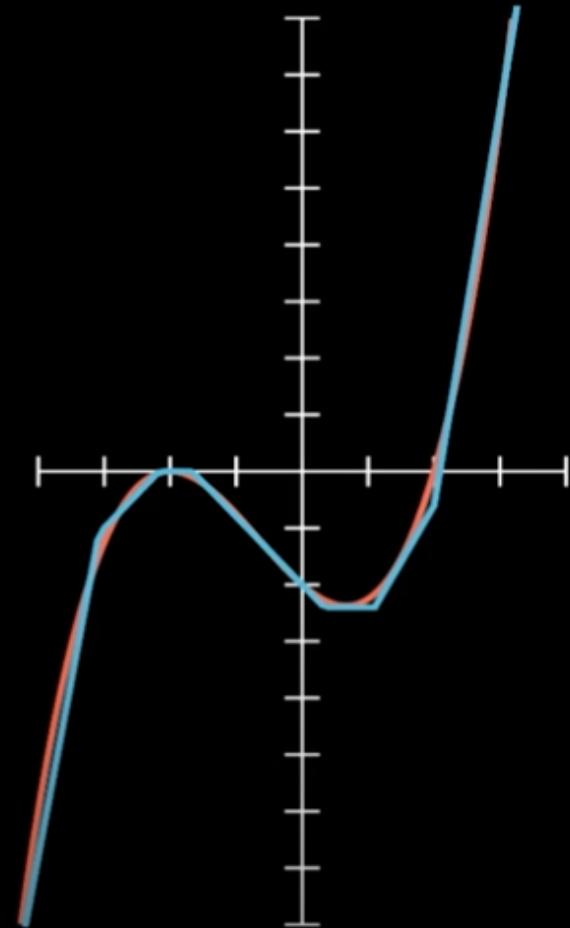


# Non-linearity



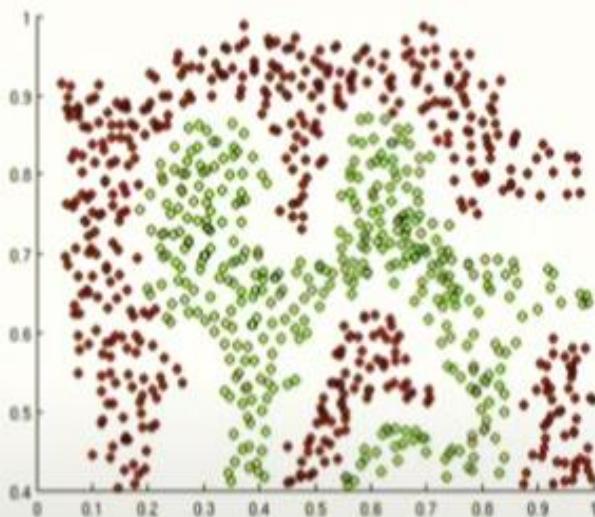
# Non-linearity

$-max(-5x + -7.7, 0) +$   
 $-max(-1.2x + -1.3, 0) +$   
 $-max(1.2x + 1, 0) +$   
 $max(1.2x + -0.2, 0) +$   
 $max(2x + -1.1, 0) +$   
 $max(5x + -5, 0)$



# Importance of Activation Functions

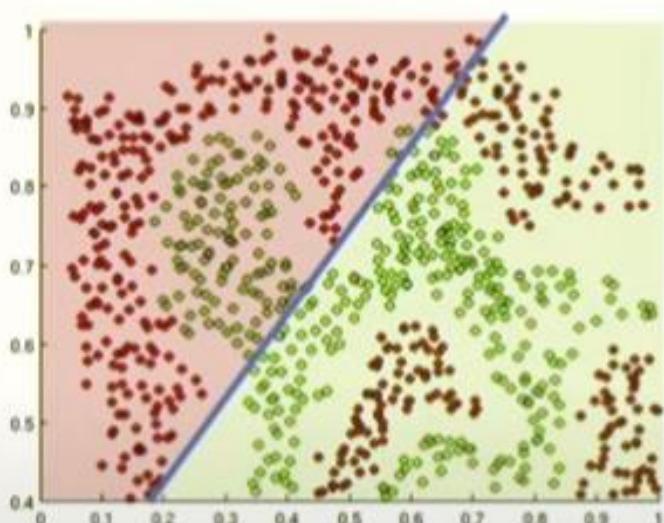
The purpose of activation functions is to *introduce non-linearities* into the network



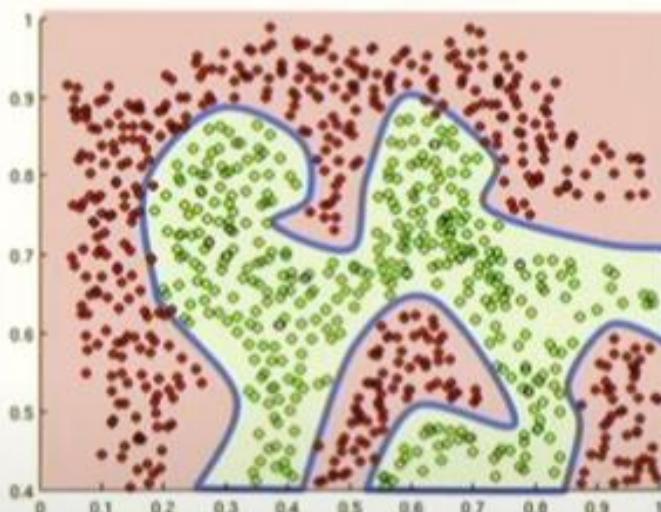
What if we wanted to build a neural network to  
distinguish green vs red points?

# Importance of Activation Functions

The purpose of activation functions is to *introduce non-linearities* into the network

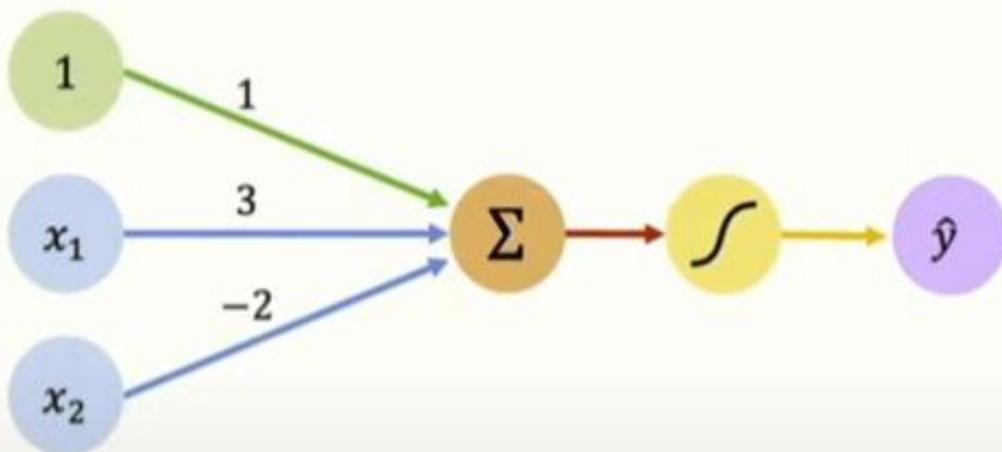


Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

# The Perceptron: Example

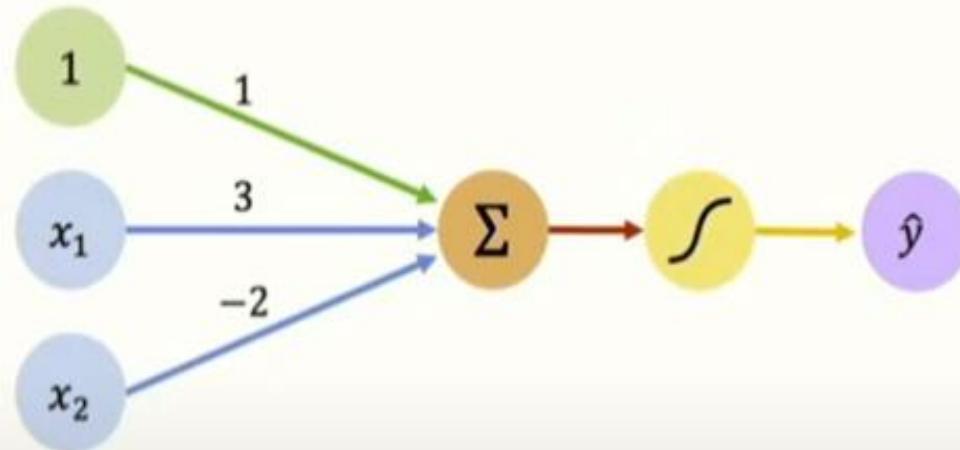


We have:  $w_0 = 1$  and  $\mathbf{w} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

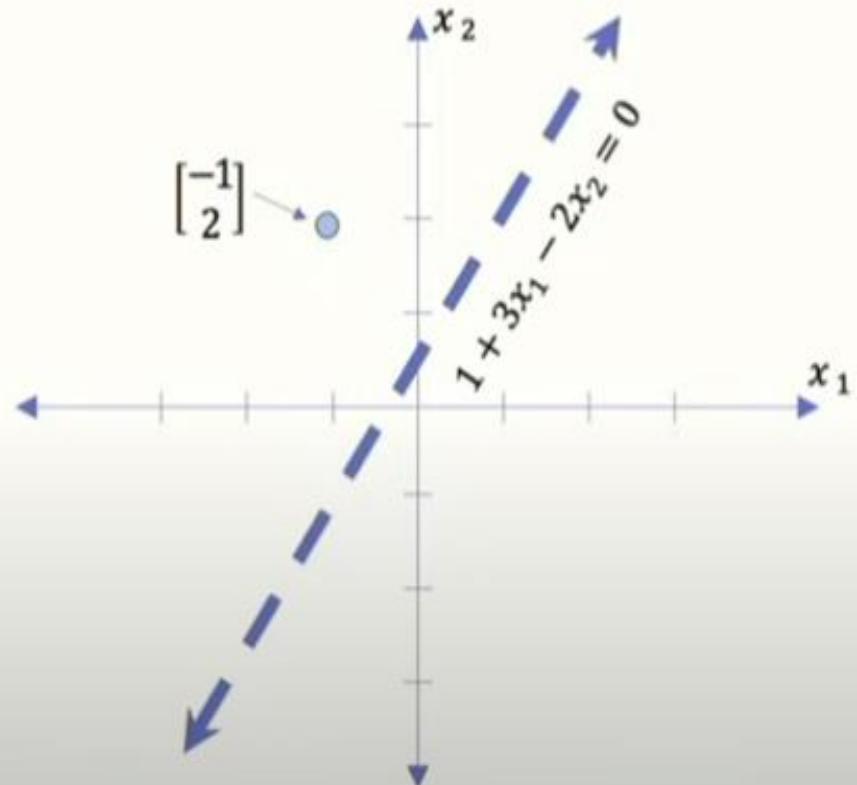
$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{w}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ \hat{y} &= g\left(1 + 3x_1 - 2x_2\right)\end{aligned}$$

This is just a line in 2D!

# The Perceptron: Example



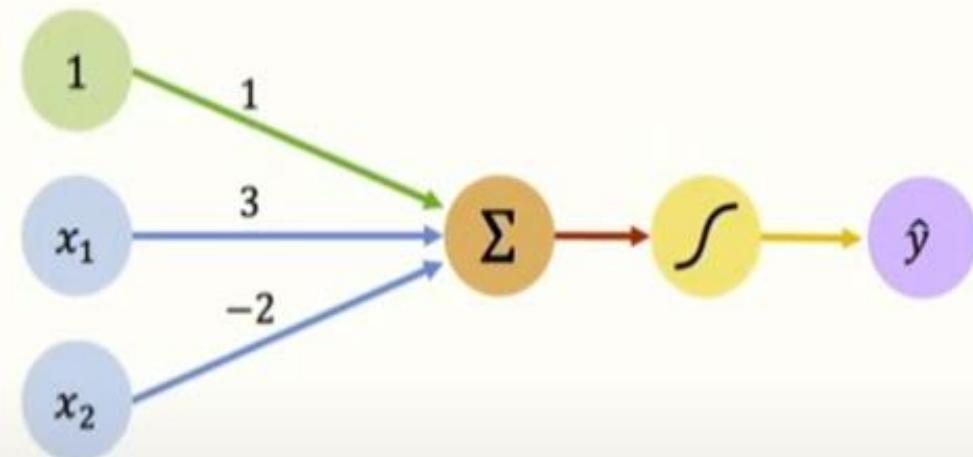
$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



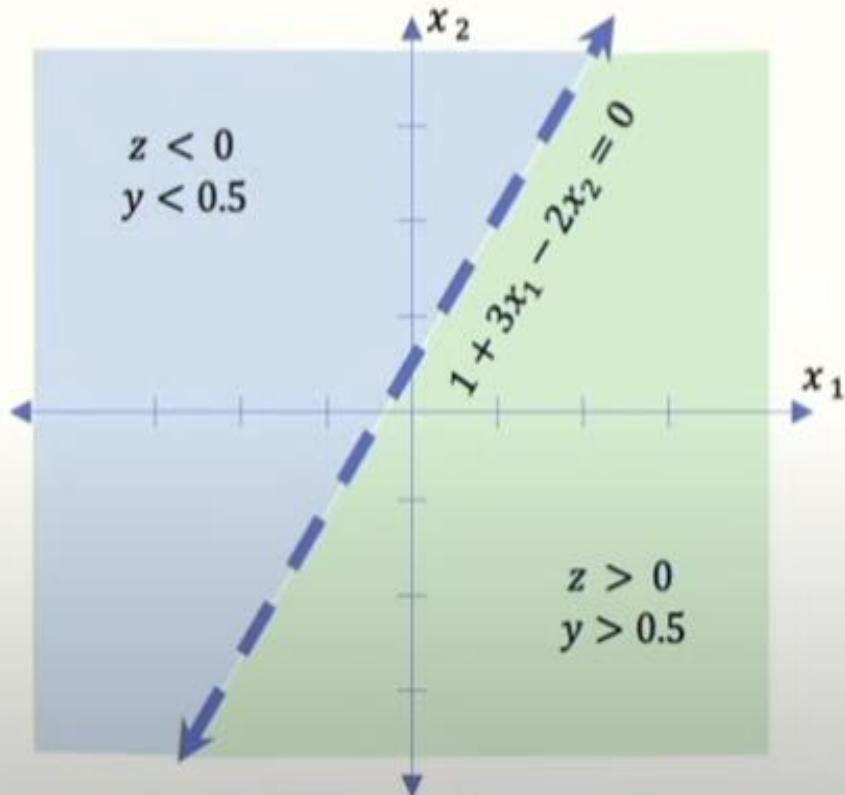
Assume we have input:  $X = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(1 + (3 * -1) - (2 * 2)) \\ &= g(-6) \approx 0.002\end{aligned}$$

# The Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



---

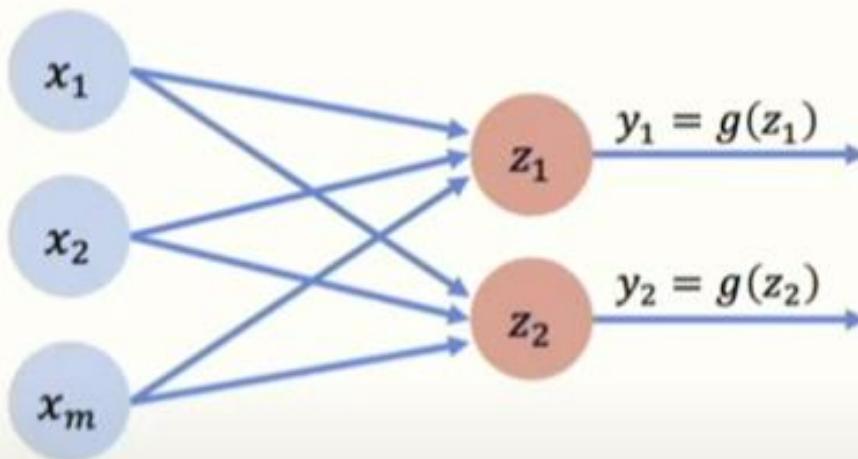
# **Building Neural Networks with Perceptrons**

# Deep Learning Memes



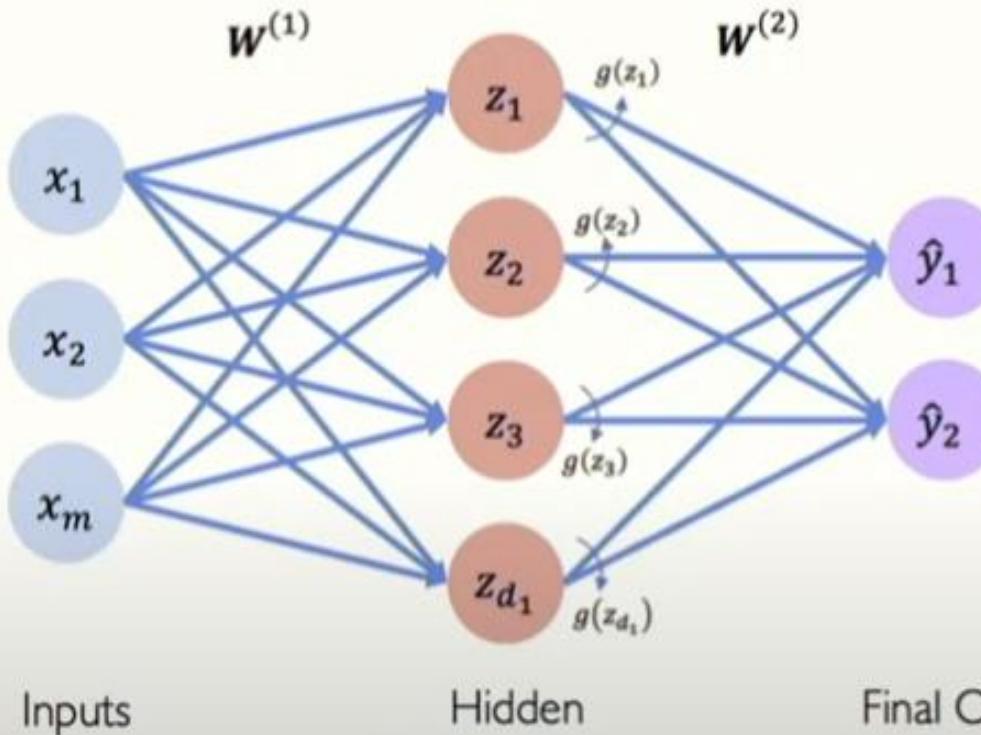
What the job of a neuron in an ANN really looks like.

# Multi Output Perceptron



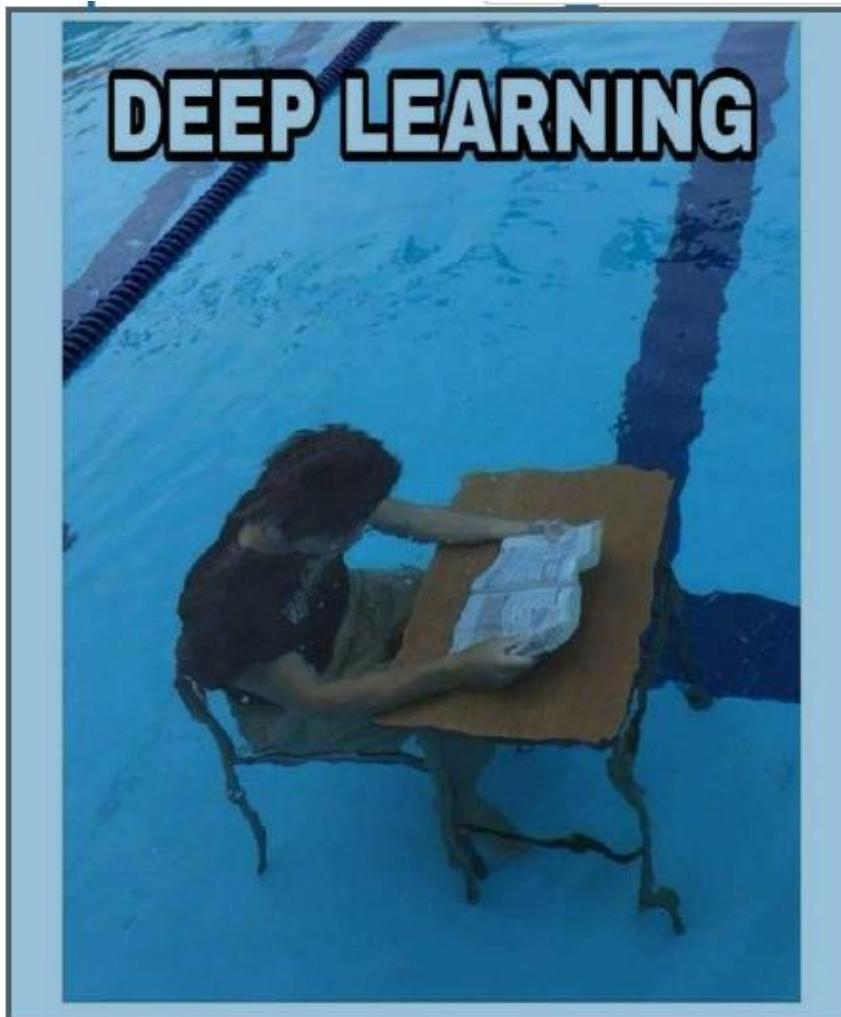
$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

# Single Layer Neural Network

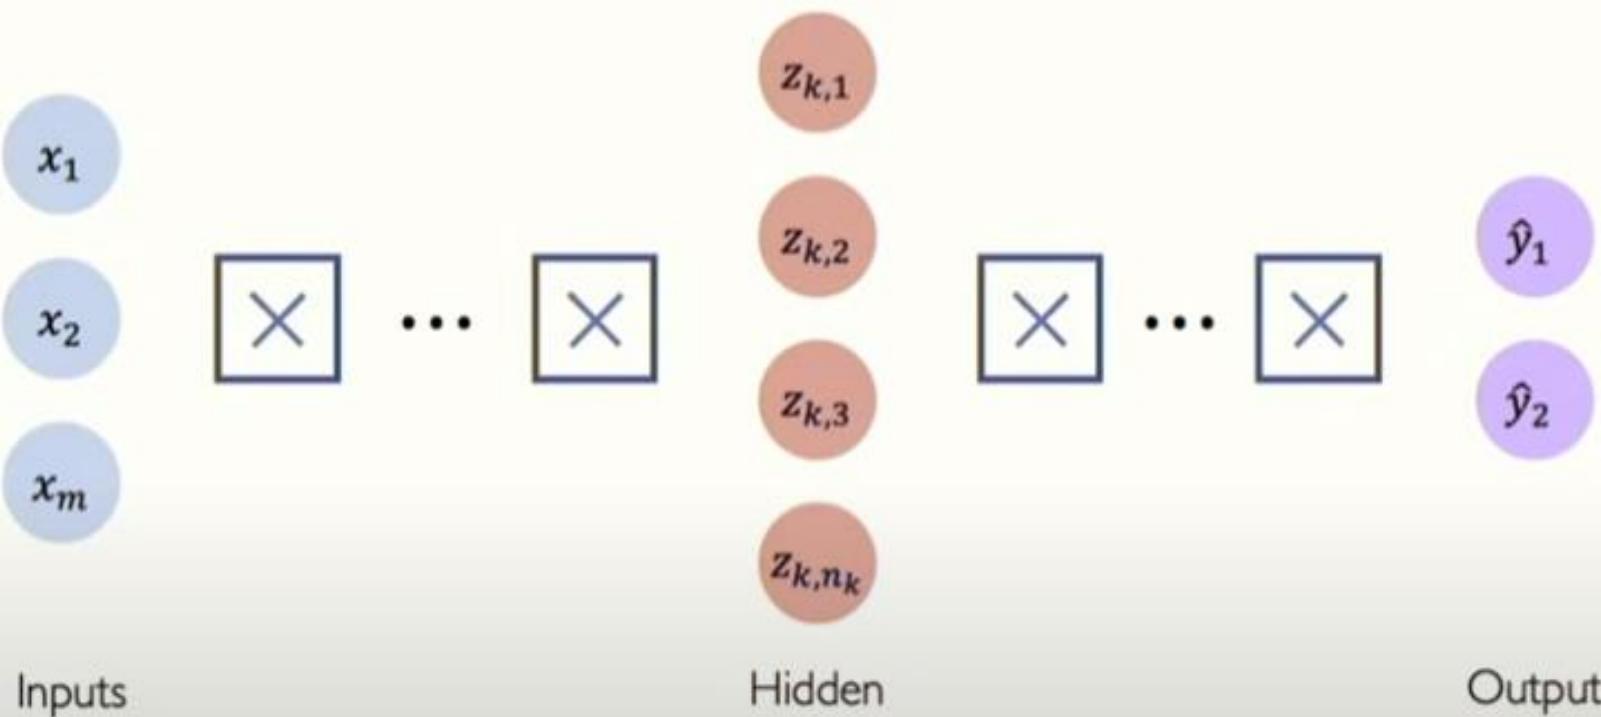


$$z_i = w_{0,i}^{(1)} + \sum_{j=1}^m x_j w_{j,i}^{(1)}$$
$$\hat{y}_i = g \left( w_{0,i}^{(2)} + \sum_{j=1}^{d_1} g(z_j) w_{j,i}^{(2)} \right)$$

# What's Deep in Deep Learning



# Deep Neural Network



$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{n_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$

# Example Problem

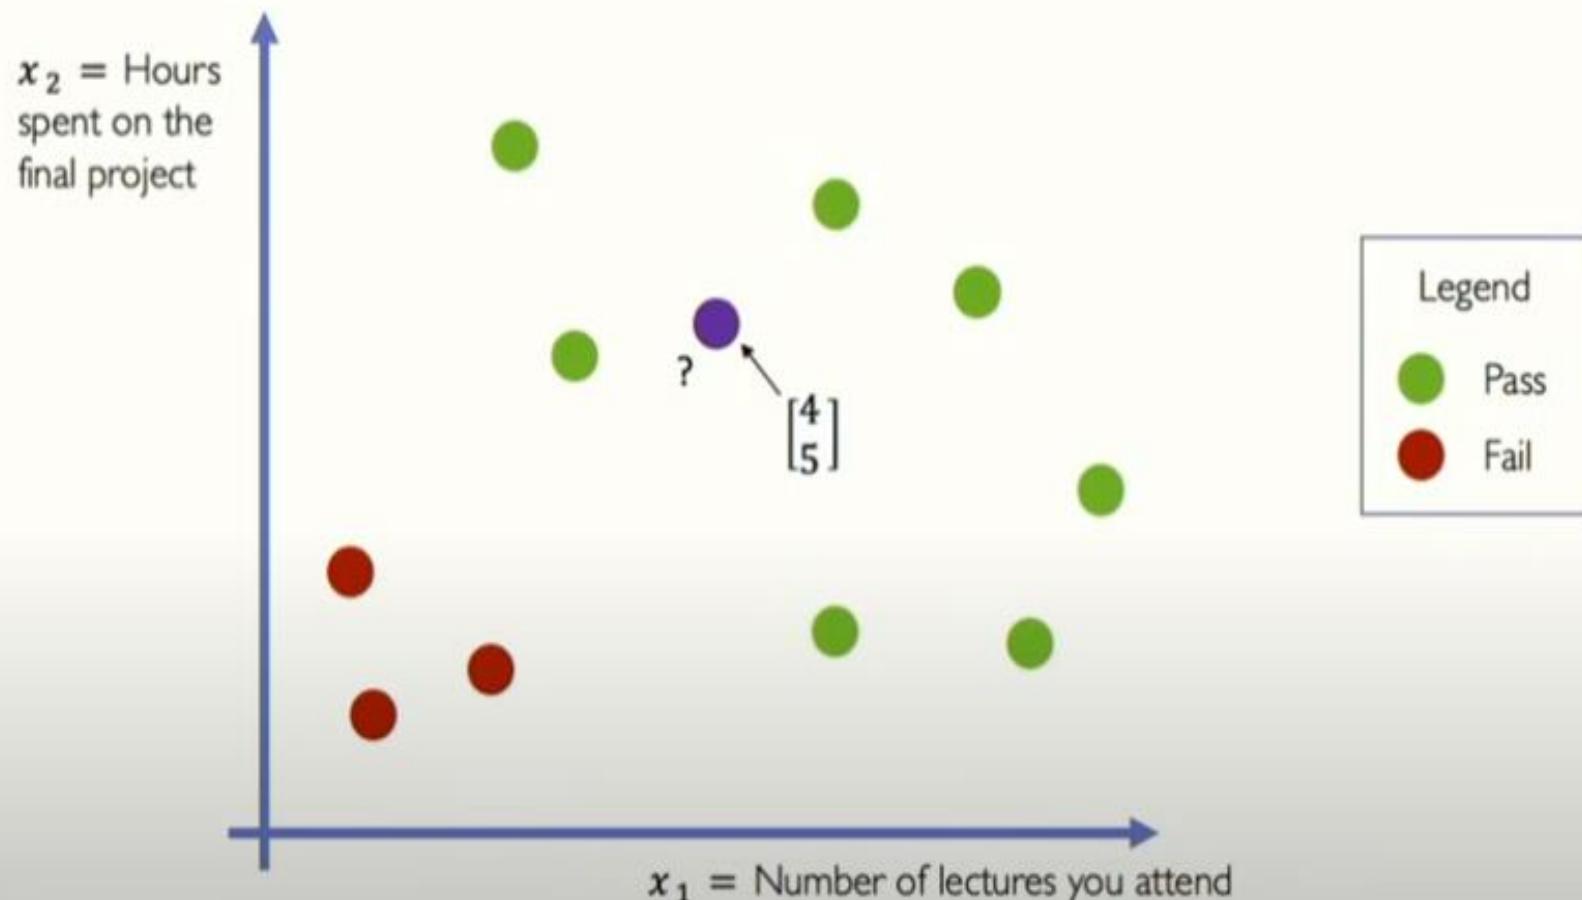
Will I pass this class?

Let's start with a simple two feature model

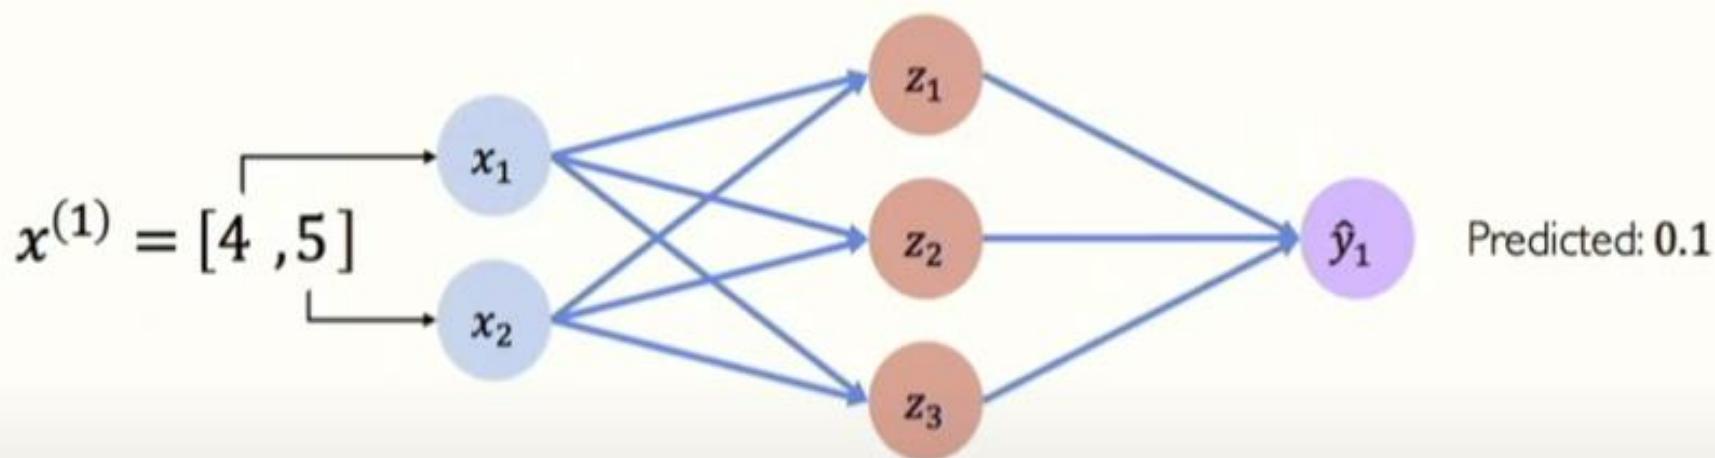
$x_1$  = Number of lectures you attend

$x_2$  = Hours spent on the final project

# Example Problem: Will I pass this class?

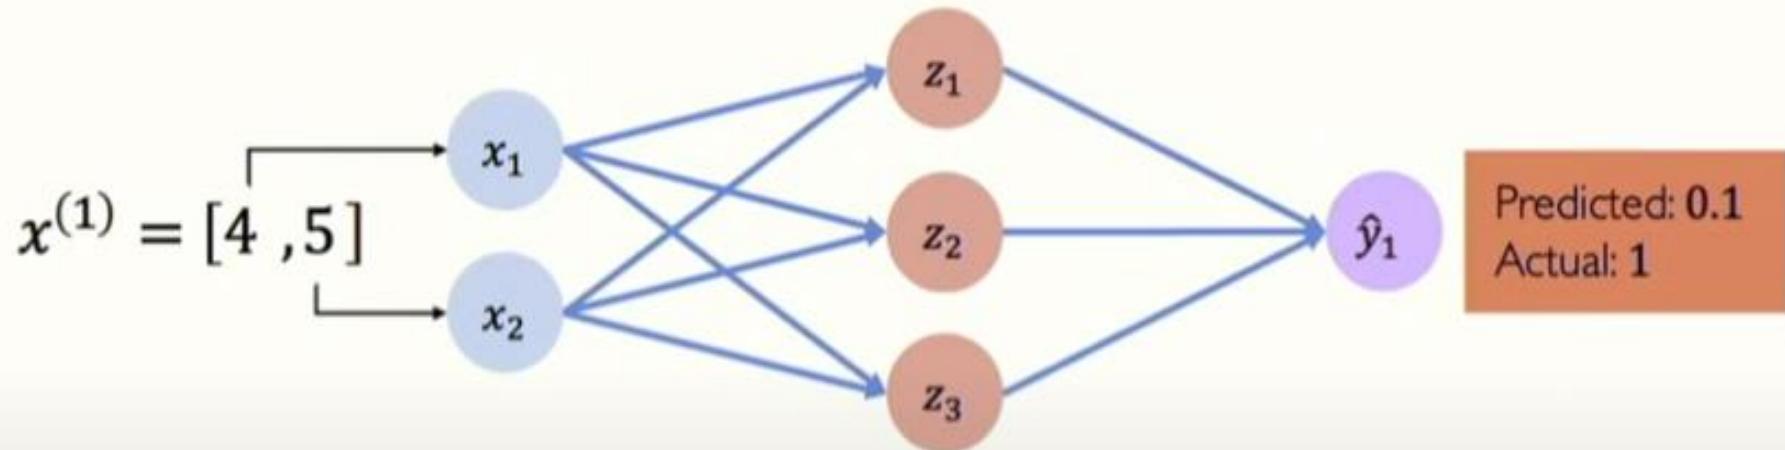


# Example Problem: Will I pass this class?



# Quantifying Loss

The **loss** of our network measures the cost incurred from incorrect predictions



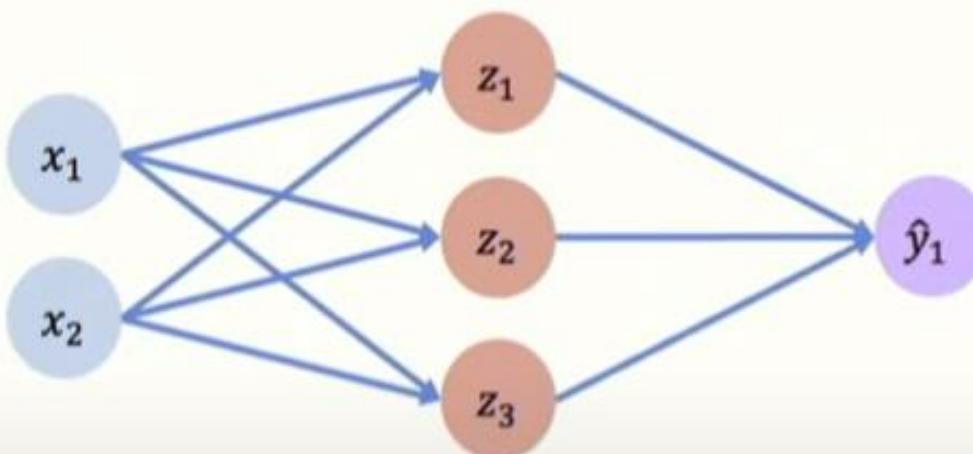
$$\mathcal{L}(f(x^{(i)}; \mathbf{w}), y^{(i)})$$

Predicted      Actual

# Empirical Loss

The *empirical loss* measures the total loss over our entire dataset

$$\mathbf{X} = \begin{bmatrix} 4, & 5 \\ 2, & 1 \\ 5, & 8 \\ \vdots & \vdots \end{bmatrix}$$



$f(x)$	$y$
0.1	✗
0.8	✗
0.6	✓
⋮	⋮

Also known as:

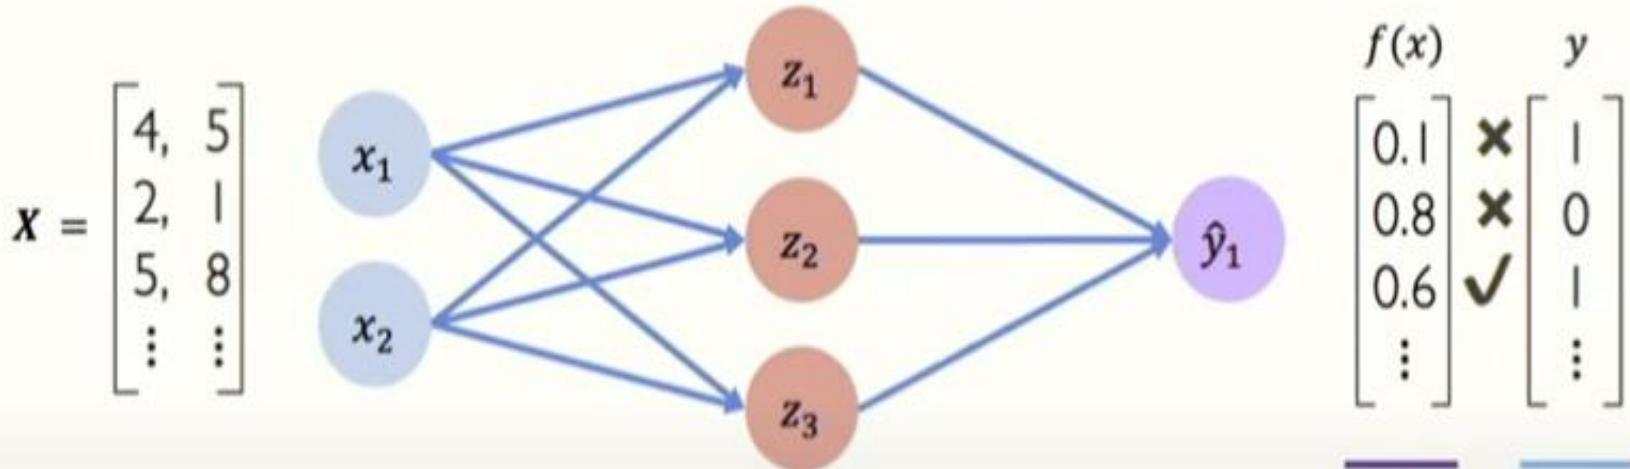
- Objective function
- Cost function
- Empirical Risk

$\mathcal{J}(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$

Predicted      Actual

# Binary Cross Entropy Loss

Cross entropy loss can be used with models that output a probability between 0 and 1

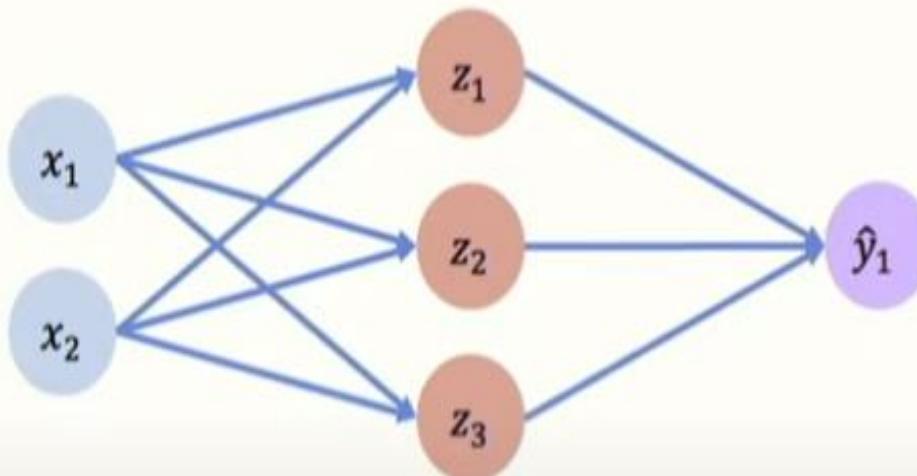


$$J(\mathbf{W}) = -\frac{1}{n} \sum_{i=1}^n \underbrace{y^{(i)} \log(f(x^{(i)}; \mathbf{W}))}_{\text{Actual}} + \underbrace{(1 - y^{(i)}) \log(1 - f(x^{(i)}; \mathbf{W}))}_{\text{Predicted}}$$

# Mean Squared Error Loss

Mean squared error loss can be used with regression models that output continuous real numbers

$$x = \begin{bmatrix} 4, & 5 \\ 2, & 1 \\ 5, & 8 \\ \vdots & \vdots \end{bmatrix}$$



$f(x)$	$y$
30	✗ 90
80	✗ 20
85	✓ 95
⋮	⋮

Final Grades  
(percentage)

$$J(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n \left( \underline{y^{(i)}} - \underline{f(x^{(i)}; \mathbf{W})} \right)^2$$

Actual      Predicted

# Loss Optimization

We want to find the network weights that **achieve the lowest loss**

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$



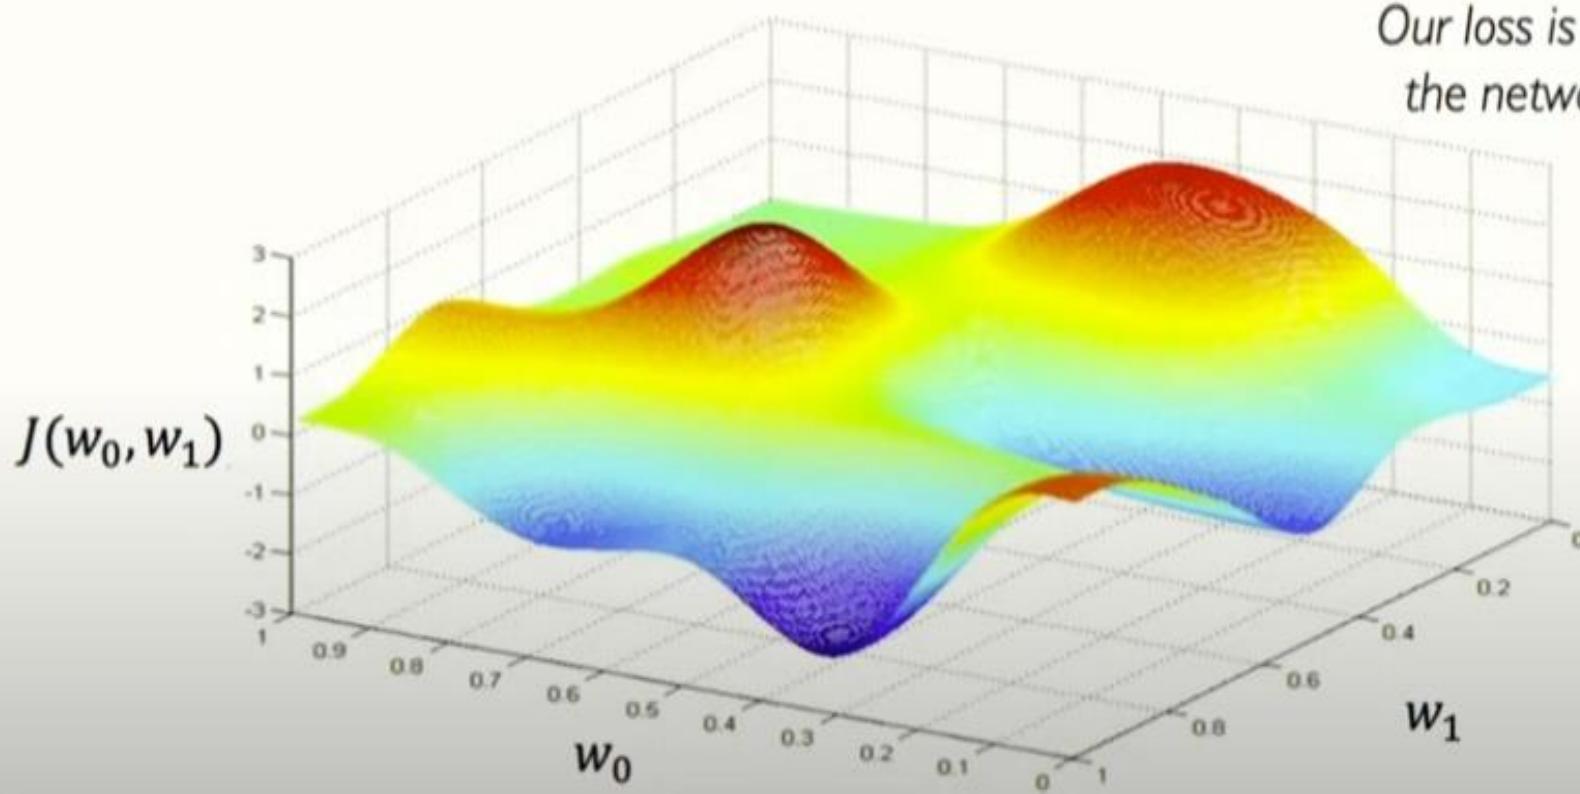
Remember:

$$\mathbf{W} = \{\mathbf{W}^{(0)}, \mathbf{W}^{(1)}, \dots\}$$

# Loss Optimization

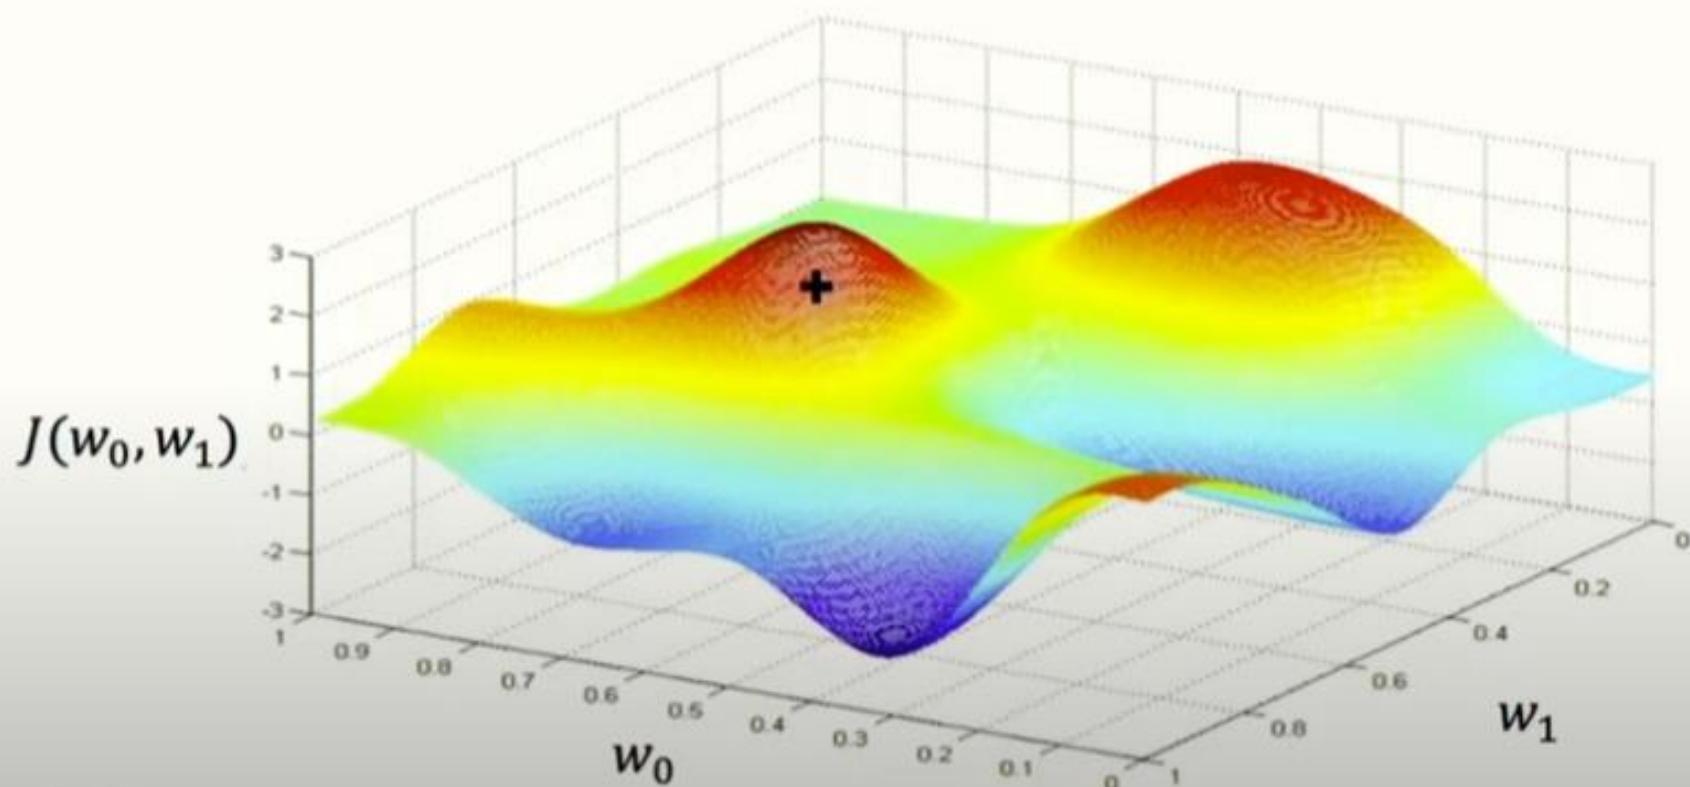
$$\mathbf{W}^* = \underset{\mathbf{W}}{\operatorname{argmin}} J(\mathbf{W})$$

Remember:  
Our loss is a function of  
the network weights!



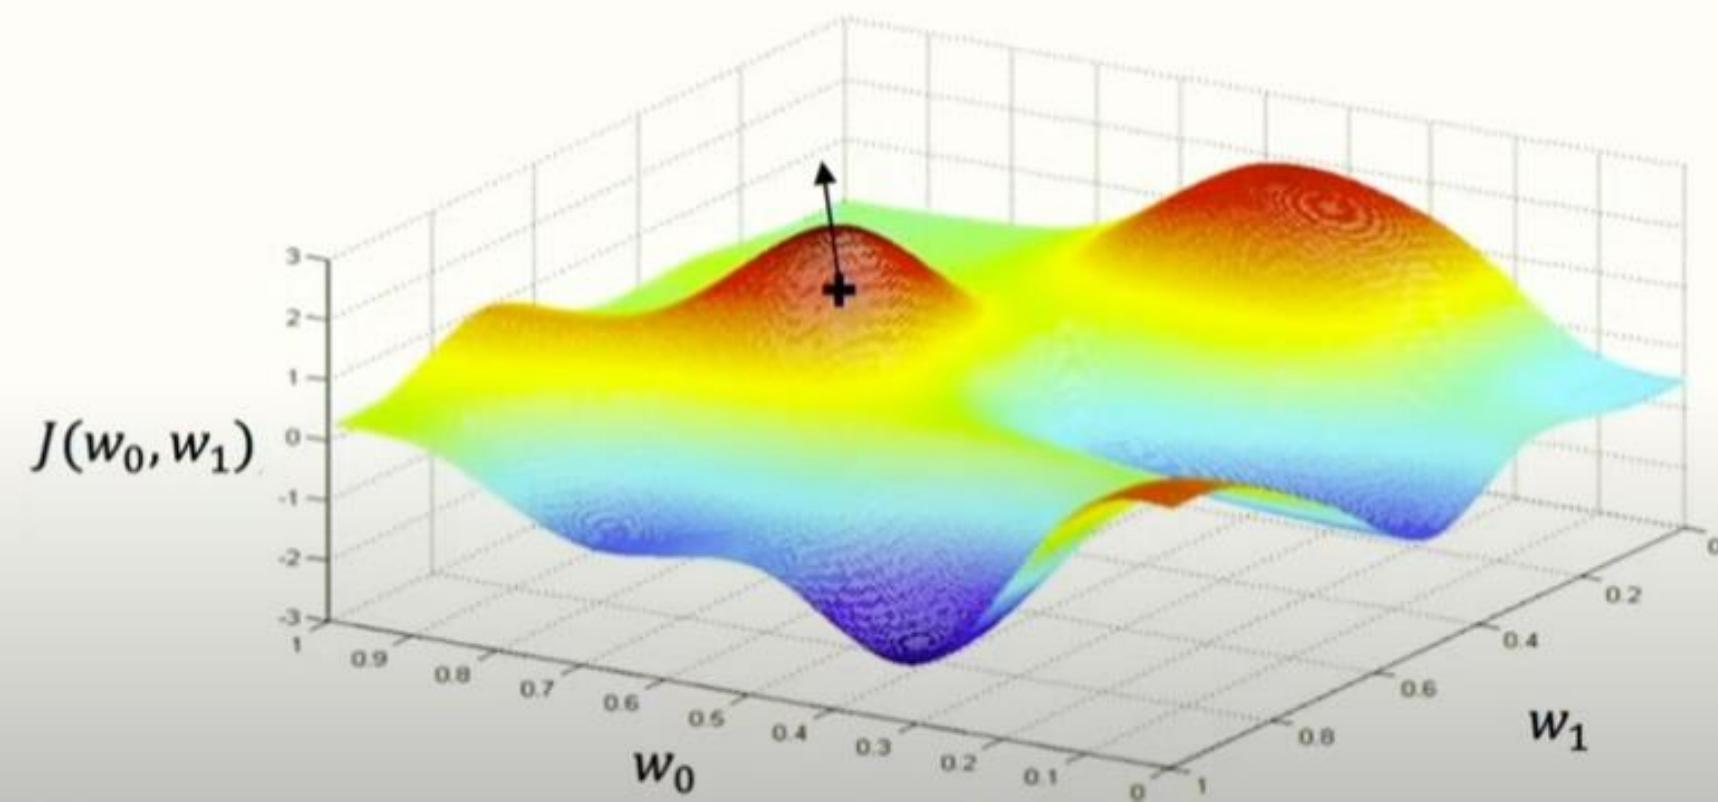
# Loss Optimization

Randomly pick an initial  $(w_0, w_1)$



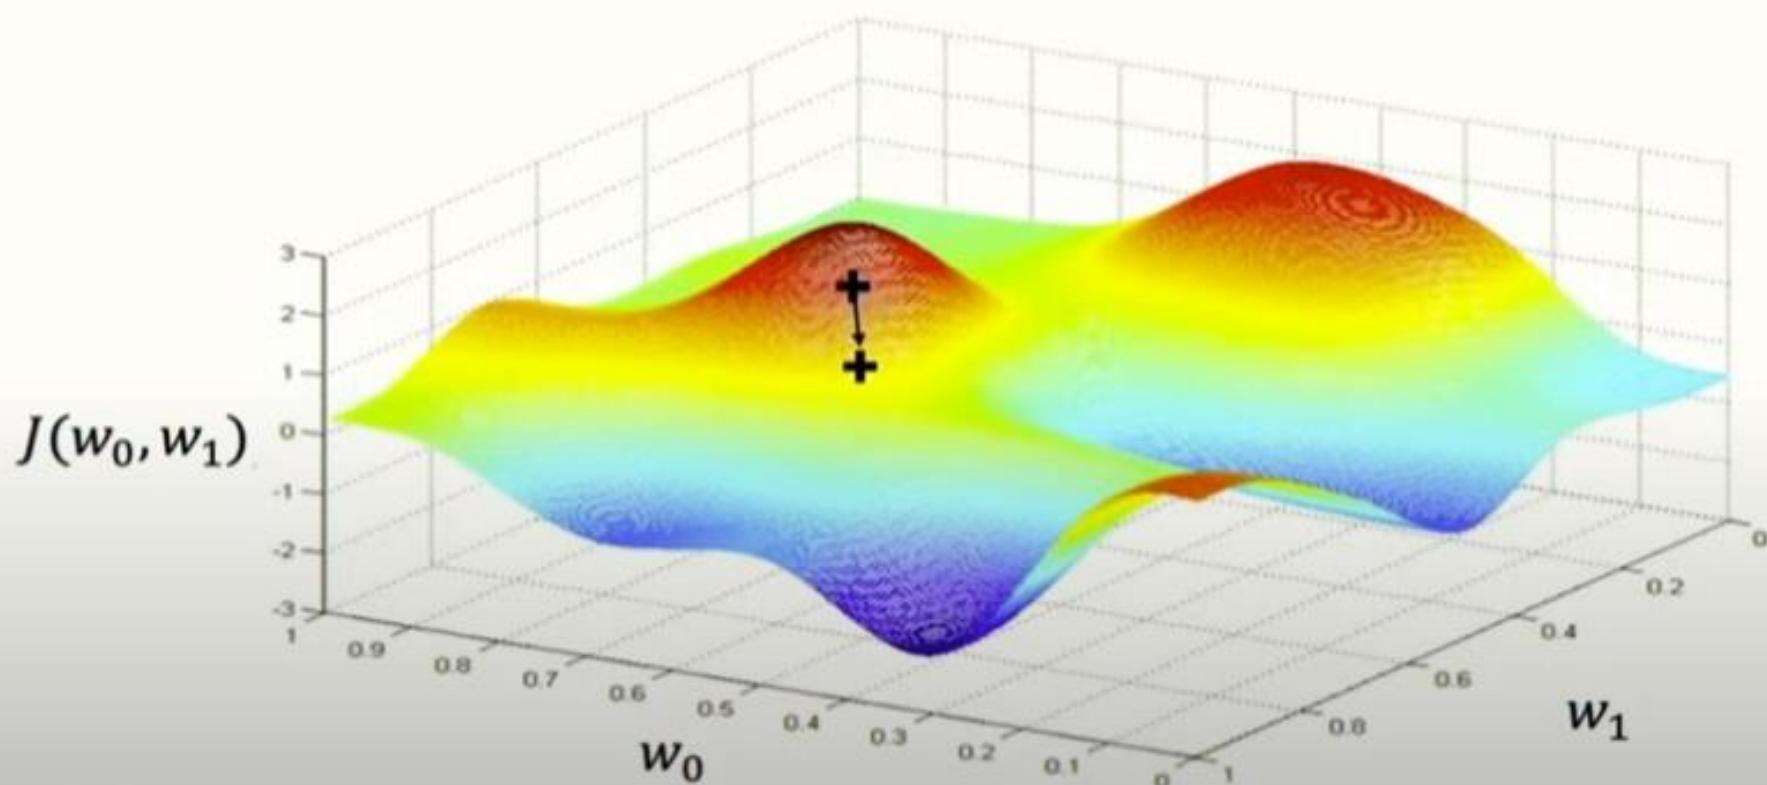
# Loss Optimization

Compute gradient,  $\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$



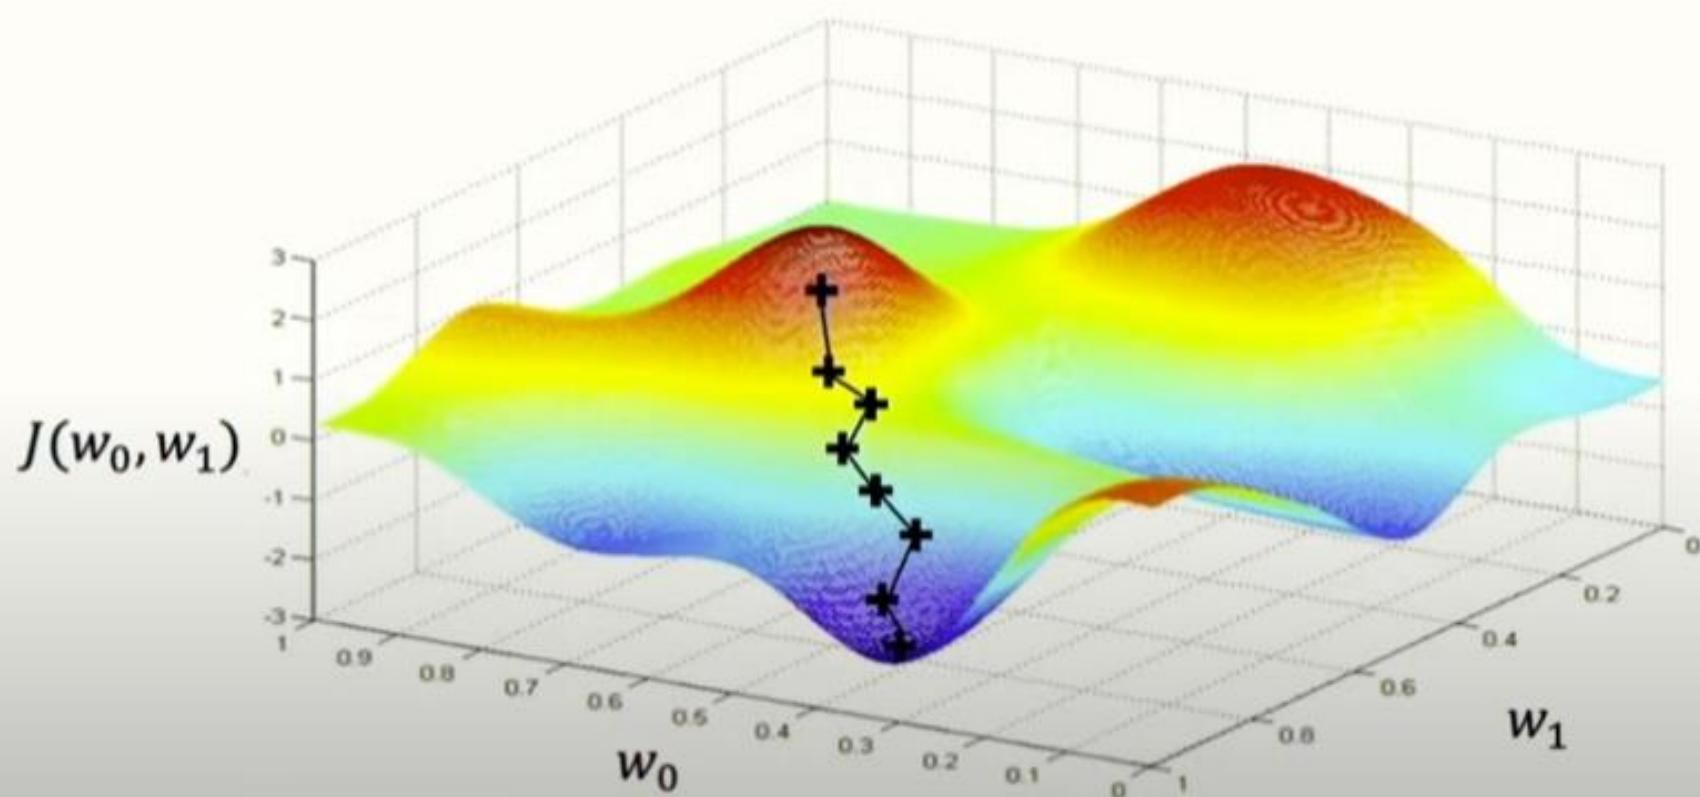
# Loss Optimization

Take small step in opposite direction of gradient



# Gradient Descent

Repeat until convergence

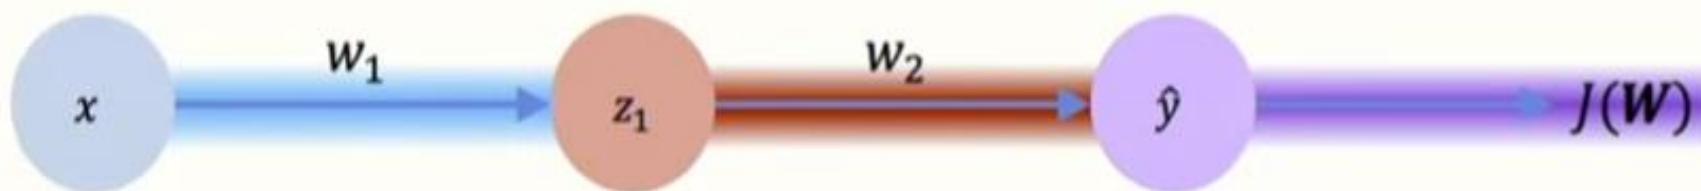


# Gradient Descent

## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient,  $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights

# Computing Gradients: Backpropagation

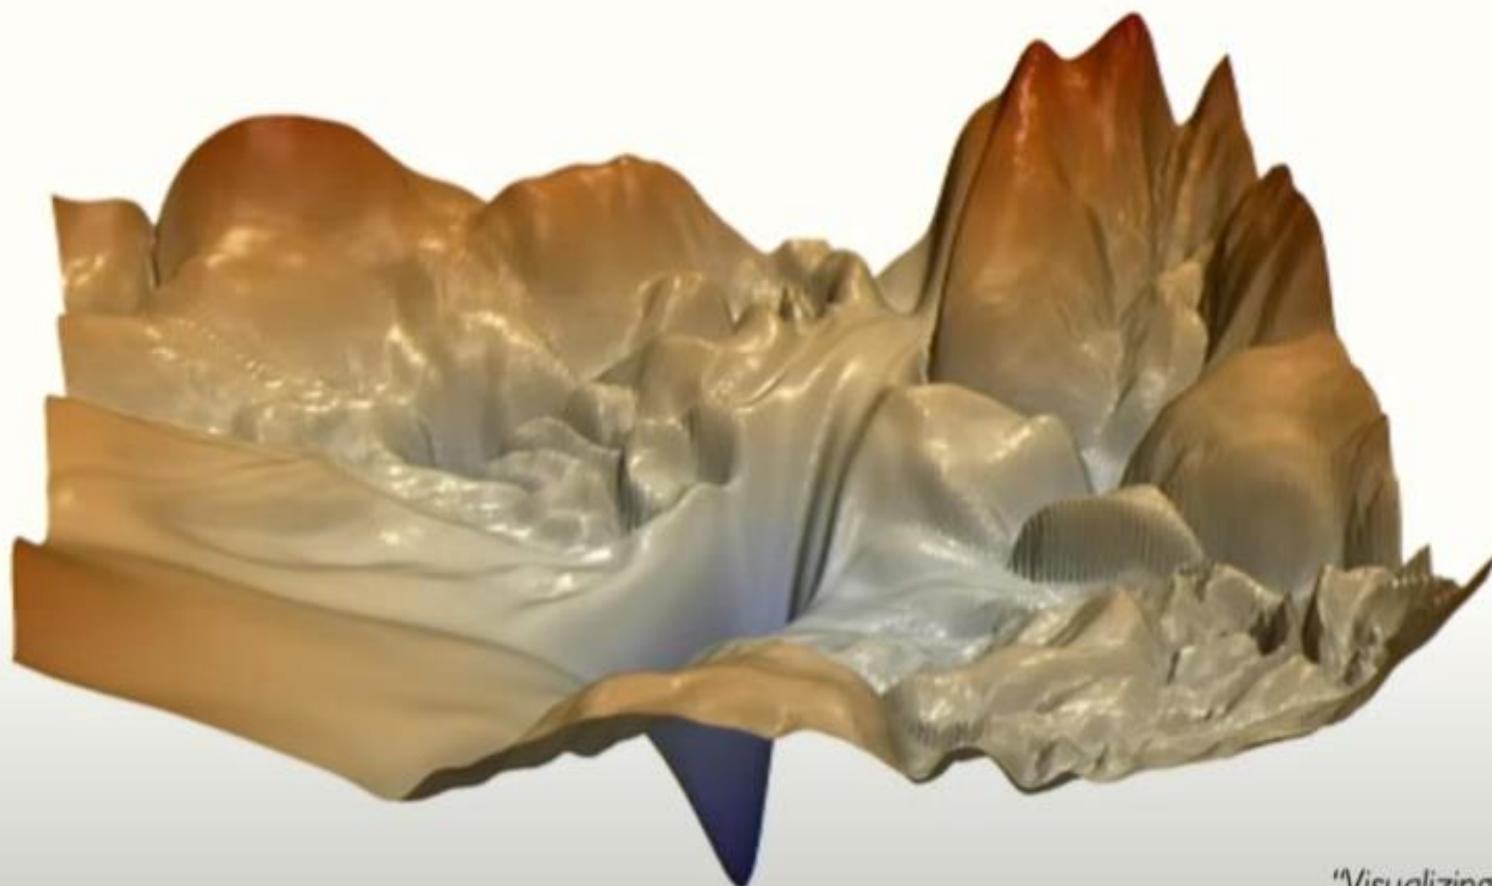


$$\frac{\partial J(\mathbf{W})}{\partial w_1} = \underline{\frac{\partial J(\mathbf{W})}{\partial \hat{y}}} * \underline{\frac{\partial \hat{y}}{\partial z_1}} * \underline{\frac{\partial z_1}{\partial w_1}}$$

— — —

Repeat this for **every weight in the network** using gradients from later layers

# Training Neural Networks is Difficult



"Visualizing the loss landscape  
of neural nets". Dec 2017.

# Loss Optimization

$$w^{(t+1)} = w^{(t)} - \alpha \nabla f(w^{(t)})$$

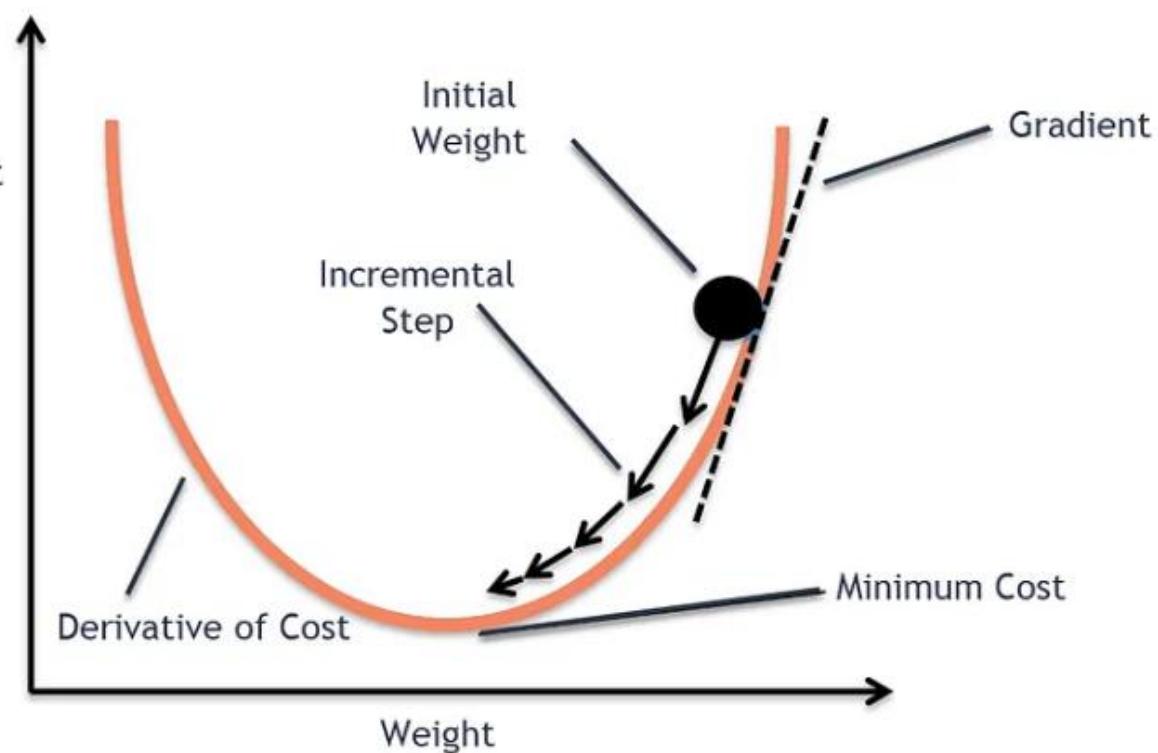
Gradient descent

position of next iteration

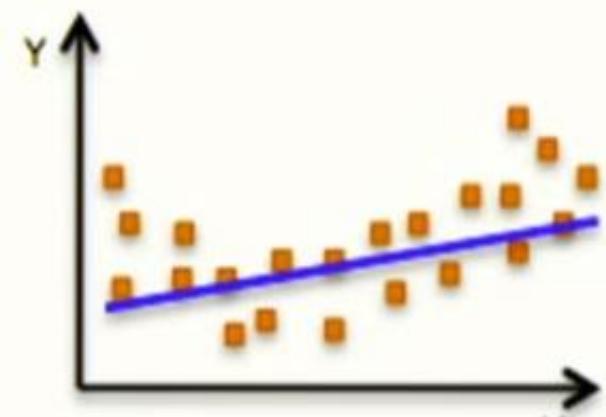
position of previous step

learning rate

step

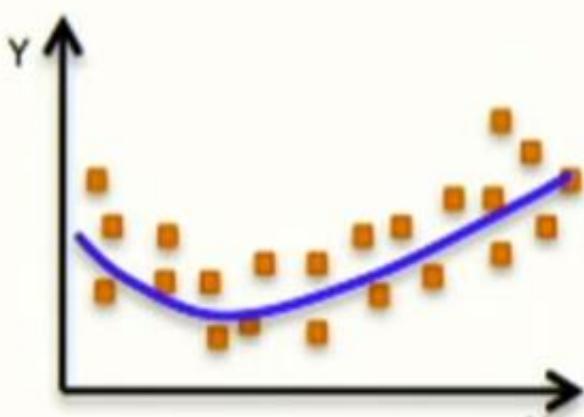


# The Problem of Overfitting

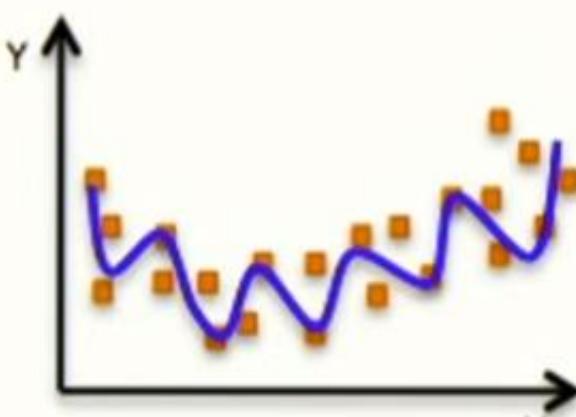


## Underfitting

Model does not have capacity  
to fully learn the data



## Ideal fit



## Overfitting

Too complex, extra parameters  
does not generalize well

# Regularization

## **What is it?**

*Technique that constrains our optimization problem to discourage complex models*

# Regularization

*What is it?*

*Technique that constrains our optimization problem to discourage complex models*

**Why do we need it?**

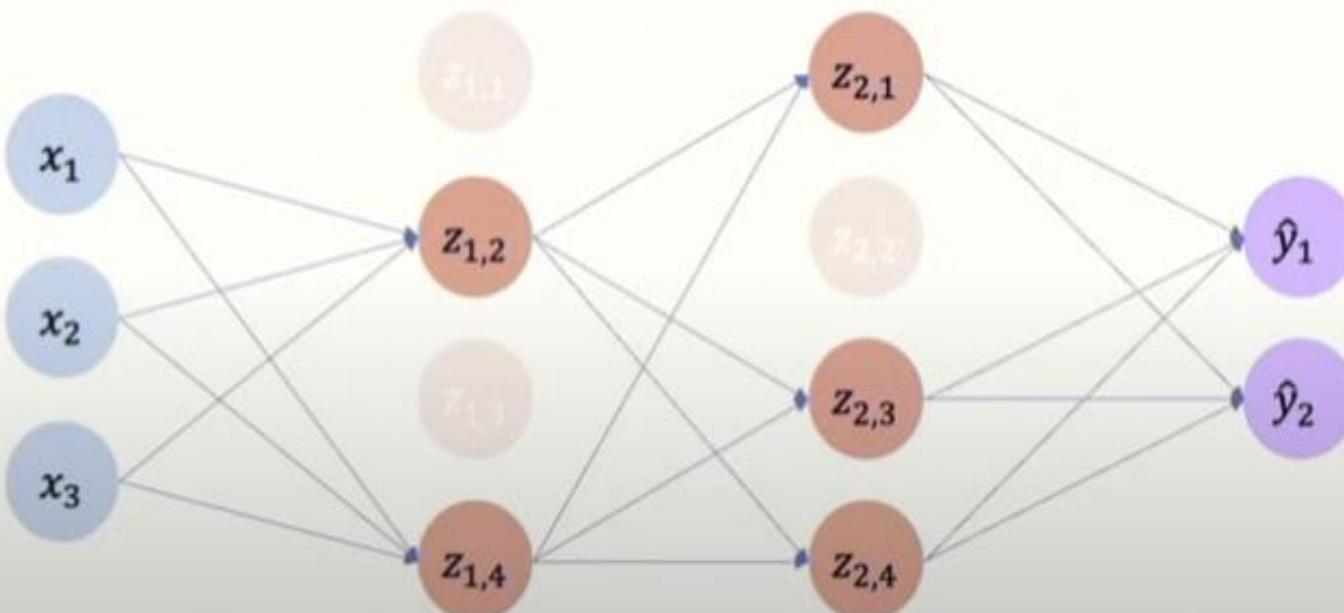
*Improve generalization of our model on unseen data*

# Regularization I: Dropout

- During training, randomly set some activations to 0
  - Typically 'drop' 50% of activations in layer
  - Forces network to not rely on any 1 node

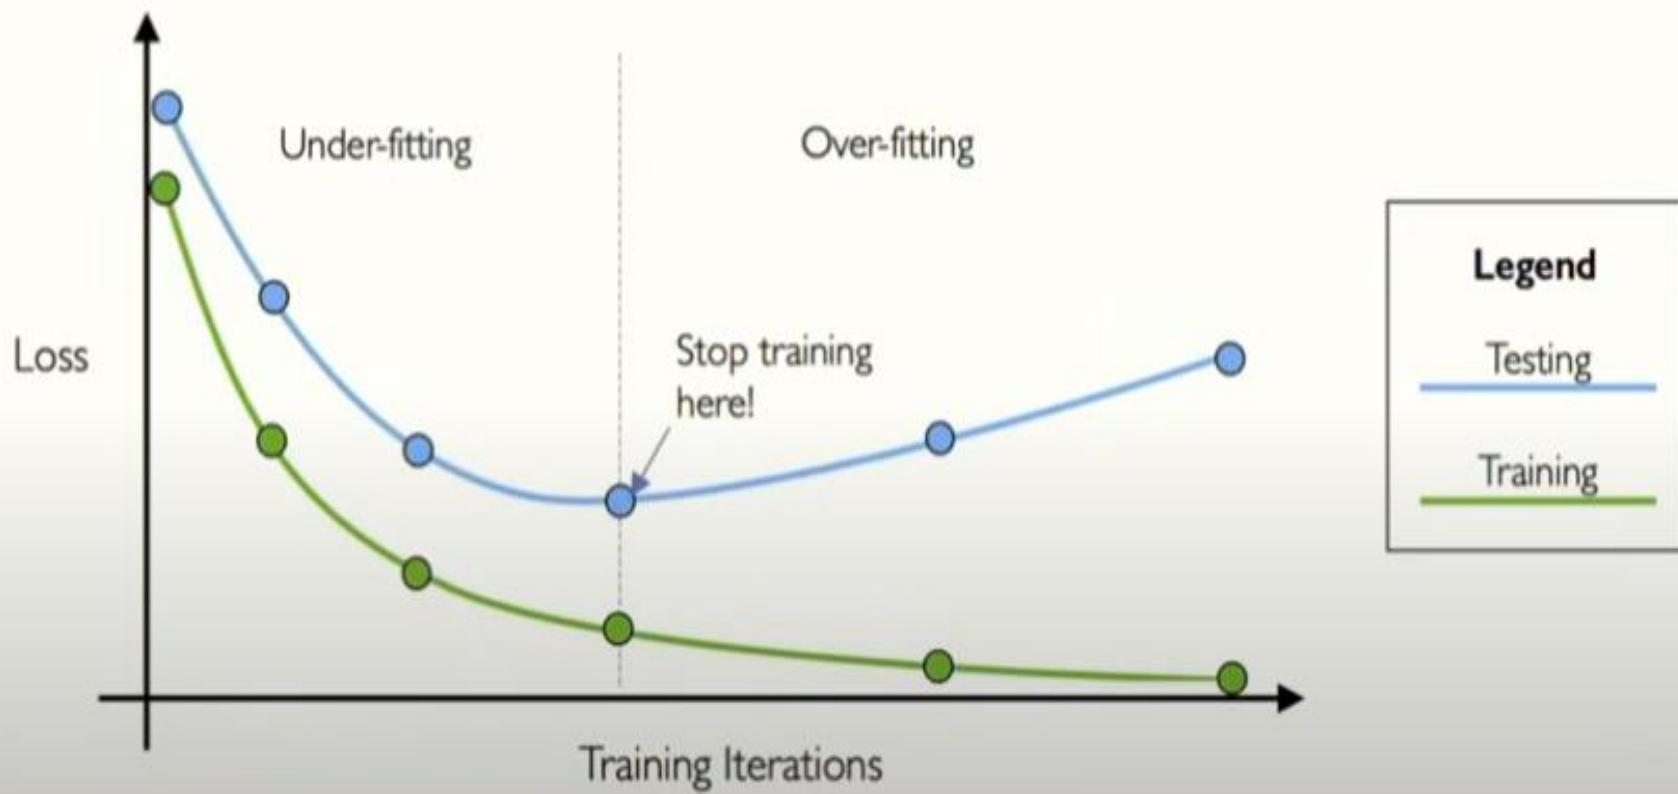


`tf.keras.layers.Dropout(p=0.5)`



# Regularization 2: Early Stopping

- Stop training before we have a chance to overfit

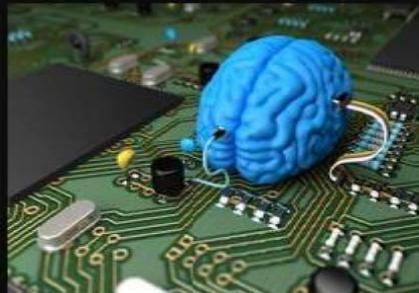


# Deep Learning Memes

## Deep Learning



What society thinks I do



What my friends think I do



What other computer  
scientists think I do



What mathematicians think I do

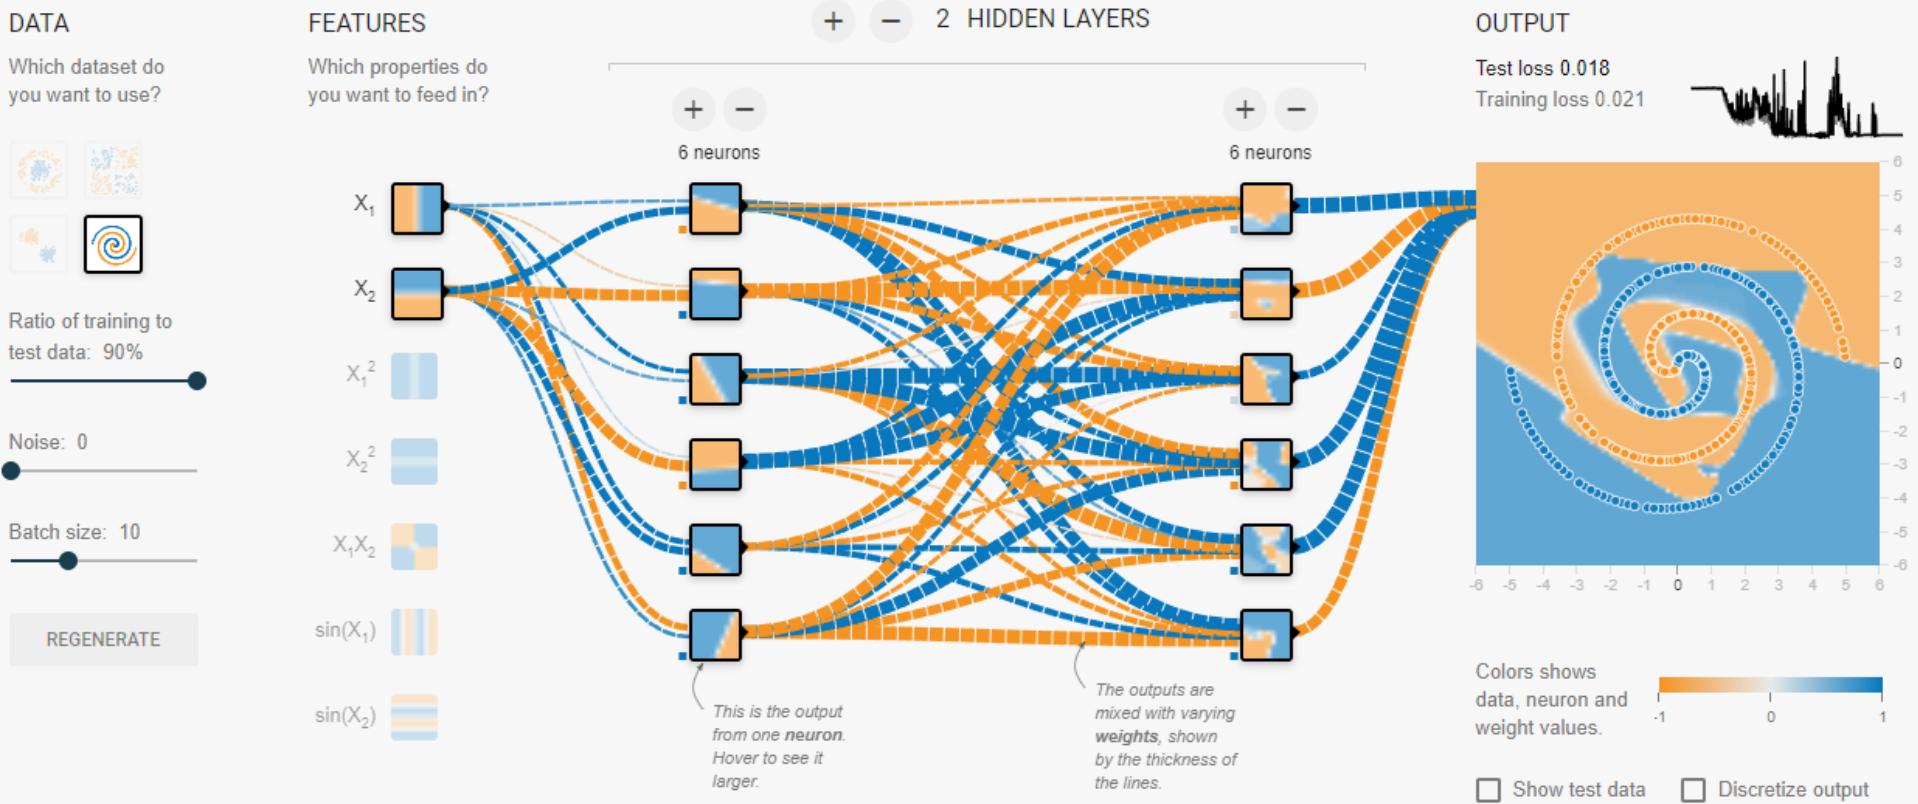


What I think I do

from theano import \*

What I actually do

# Fully Connect Artifical Neural Network



---

**Thank You.**